

BUSINESS UNDERSTANDING

This project analyzes movie performance data to uncover the key factors that drive box office success and audience ratings. By examining variables such as genre, budget, runtime, ratings, and contributor roles (directors and writers), it aims to help studios make data-driven decisions in budgeting, hiring, and marketing. The goal is to identify high-impact talent, understand how production choices affect profitability, and provide actionable insights to improve investment, content, and strategic planning in the film industry.

SECTION 1: DATA UNDERSTANDING

The datasets used in this project are sourced from publicly available movie databases, including Box Office Mojo (BOM), The Numbers (TN), IMDB database and The Movie Database (TMDb). They provide structured information on thousands of films released over the past two decades. The data spans multiple formats; categorical variables such as genre, director, and writer, numerical variables such as budget, gross revenue, and ratings and temporal variables such as release year and runtime. The dataset includes:

- Film titles & release dates
- Genres (one or multiple per film)
- Production budget & box office revenue
- Ratings & popularity metrics
- Contributor data (directors, writers etc.)
- Movie runtime

```
In [1]: # imports
import pandas as pd # cleaning data
import matplotlib.pyplot as plt # data visualization
import seaborn as sns # data visualization
import numpy as np # python calculations
import json # parse data
import warnings # ignore warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
In [2]: # import all required datasets
length_data = pd.read_csv("Data/rt.movie_info.tsv", sep='\t')
ratings_data = pd.read_csv("Data/tmdb.movies.csv")
merged_movie_info = pd.read_csv("Data/merged_movie_info.csv", index_col=0)
df = pd.read_csv('Data/clean_full_movie_data.csv')
```

In [3]: # check the first rows of dataset
length_data.head()

Out[3]:

| | id | synopsis | rating | genre | director | writer | theater_date |
|---|-----------|---|---------------|-------------------------------------|------------------|---------------------------------|---------------------|
| 0 | 1 | This gritty, fast-paced, and innovative police... | R | Action and Adventure Classics Drama | William Friedkin | Ernest Tidyman | Oct 9, 1971 |
| 1 | 3 | New York City, not-too-distant-future: Eric Pa... | R | Drama Science Fiction and Fantasy | David Cronenberg | David Cronenberg Don DeLillo | Aug 17, 2012 |
| 2 | 5 | Illeana Douglas delivers a superb performance ... | R | Drama Musical and Performing Arts | Allison Anders | Allison Anders | Sep 13, 1996 |
| 3 | 6 | Michael Douglas runs afoul of a treacherous su... | R | Drama Mystery and Suspense | Barry Levinson | Paul Attanasio Michael Crichton | Dec 9, 1994 |
| 4 | 7 | NaN | NR | Drama Romance | Rodney Bennett | Giles Cooper | NaN |

In [4]: # check first rows of dataset
ratings_data.head()

Out[4]:

| | Unnamed: 0 | genre_ids | id | original_language | original_title | popularity | release_date | |
|---|---------------|---------------------|-------|-------------------|--|------------|--------------|--------|
| 0 | 0 | [12, 14, 10751] | 12444 | en | Harry Potter and the Deathly Hallows: Part 1 | 33.533 | 2010-11-19 | a D Hæ |
| 1 | 1 | [14, 12, 16, 10751] | 10191 | en | How to Train Your Dragon | 28.734 | 2010-03-26 | C |
| 2 | 2 | [12, 28, 878] | 10138 | en | Iron Man 2 | 28.515 | 2010-05-07 | Iro |
| 3 | 3 | [16, 35, 10751] | 862 | en | Toy Story | 28.005 | 1995-11-22 | |
| 4 | 4 | [28, 878, 12] | 27205 | en | Inception | 27.920 | 2010-07-16 | Inc |

In [5]: # check dataset information

```
length_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1560 entries, 0 to 1559
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               1560 non-null    int64  
 1   synopsis         1498 non-null    object  
 2   rating            1557 non-null    object  
 3   genre             1552 non-null    object  
 4   director          1361 non-null    object  
 5   writer            1111 non-null    object  
 6   theater_date     1201 non-null    object  
 7   dvd_date          1201 non-null    object  
 8   currency          340 non-null    object  
 9   box_office        340 non-null    object  
 10  runtime           1530 non-null    object  
 11  studio            494 non-null    object  
dtypes: int64(1), object(11)
memory usage: 146.4+ KB
```

In [6]: # check dataset information

```
ratings_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26517 entries, 0 to 26516
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        26517 non-null    int64  
 1   genre_ids         26517 non-null    object  
 2   id                26517 non-null    int64  
 3   original_language 26517 non-null    object  
 4   original_title    26517 non-null    object  
 5   popularity         26517 non-null    float64 
 6   release_date      26517 non-null    object  
 7   title              26517 non-null    object  
 8   vote_average       26517 non-null    float64 
 9   vote_count         26517 non-null    int64  
dtypes: float64(2), int64(3), object(5)
memory usage: 2.0+ MB
```

In [7]: # look at the first five rows
merged_movie_info.head()

Out[7]:

| | | title | studio | year | production_budget | worldwide_gross | popularity | vote_average |
|----|---|-------|--------|------|-------------------|-----------------|------------|--------------|
| 1 | Inception | WB | 2010 | | 160000000 | 835524642 | 27.920 | 8.3 |
| 8 | The Chronicles of Narnia: The Voyage of the Da... | Fox | 2010 | | 155000000 | 418186950 | 17.382 | 6.3 |
| 30 | Gulliver's Travels | Fox | 2010 | | 112000000 | 232017848 | 10.768 | 5.1 |
| 34 | Due Date | WB | 2010 | | 65000000 | 211739043 | 12.445 | 6.3 |
| 36 | Yogi Bear | WB | 2010 | | 80000000 | 204774690 | 9.096 | 5.3 |

In [8]: # checkout dataset information
merged_movie_info.info()

```
<class 'pandas.core.frame.DataFrame'>
Index: 282 entries, 1 to 2470
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   title            282 non-null    object 
 1   studio           282 non-null    object 
 2   year             282 non-null    int64  
 3   production_budget 282 non-null    int64  
 4   worldwide_gross  282 non-null    int64  
 5   popularity        282 non-null    float64
 6   vote_average      282 non-null    float64
 7   vote_count        282 non-null    int64  
 8   profit            282 non-null    int64  
 9   genre             282 non-null    object 
 10  director          282 non-null    object 
 11  writer            282 non-null    object 
dtypes: float64(2), int64(5), object(5)
memory usage: 28.6+ KB
```

In [9]: # check the first rows
df.head()

Out[9]:

| | movie_titles | genres | original_language | release_date | studio | production_budget | dom |
|---|----------------------------|--|-------------------|--------------|--------|-------------------|-----|
| 0 | toy story 3 | [{"id": 16, "name": "Animation"}, {"id": 10751...} | en | 2010-06-16 | BV | 2000000000 | \$ |
| 1 | inception | [{"id": 28, "name": "Action"}, {"id": 53, "nam...} | en | 2010-07-14 | WB | 1600000000 | \$ |
| 2 | shrek forever after | [{"id": 35, "name": "Comedy"}, {"id": 12, "nam...} | en | 2010-05-16 | P/DW | 1650000000 | \$ |
| 3 | the twilight saga: eclipse | [{"id": 12, "name": "Adventure"}, {"id": 14, "...} | en | 2010-06-23 | Sum. | 680000000 | \$ |
| 4 | iron man 2 | [{"id": 12, "name": "Adventure"}, {"id": 28, "...} | en | 2010-04-28 | Par. | 2000000000 | \$ |



In [10]: # check dataset information
df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45749 entries, 0 to 45748
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   movie_titles    45748 non-null   object 
 1   genres          45603 non-null   object 
 2   original_language 45592 non-null   object 
 3   release_date    45748 non-null   object 
 4   studio          1469 non-null    object 
 5   production_budget 45748 non-null   object 
 6   domestic_gross   1470 non-null    object 
 7   foreign_gross    1274 non-null    object 
 8   worldwide_gross  1470 non-null    object 
 9   revenue          45748 non-null   object 
 10  runtime          45355 non-null   float64
 11  popularity       45603 non-null   float64
 12  production_companies 45603 non-null   object 
 13  production_countries 45603 non-null   object 
 14  vote_average     45603 non-null   float64
 15  vote_count       45603 non-null   float64
dtypes: float64(4), object(12)
memory usage: 5.6+ MB
```

SECTION 2: DATA PREPARATION

To ensure reliable analysis, the following steps were performed:

- Merging of the datasets
- Cleaned inconsistent formatting and missing values
- Converted monetary values to numeric types
- Extracted and separated multiple genres per film
- Calculated profit and categorized release dates into seasons

```
In [11]: ┌ #merging the ratings and length datasets
length_rating_df = pd.merge(
    length_data[["id", "runtime"]],
    ratings_data[["id", "vote_average", "vote_count"]],
    on="id",
    how="inner"
)
length_rating_df.head()
```

Out[11]:

| | id | runtime | vote_average | vote_count |
|----------|-----------|----------------|---------------------|-------------------|
| 0 | 27 | NaN | 4.9 | 170 |
| 1 | 90 | 96 minutes | 7.1 | 1827 |
| 2 | 93 | 110 minutes | 7.9 | 359 |
| 3 | 95 | 116 minutes | 6.7 | 4267 |
| 4 | 189 | 94 minutes | 6.3 | 2210 |

```
In [12]: ┌ # converting the columns to numeric values
length_rating_df["runtime"] = length_rating_df["runtime"].astype(str)
length_rating_df["runtime"] = pd.to_numeric(length_rating_df["runtime"])
```

```
In [13]: ┌ # check duplicated values
print(length_rating_df.duplicated().value_counts())
```

```
False    27
True     5
Name: count, dtype: int64
```

In [14]: ➔ `#drop duplicates`
`length_rating_df = length_rating_df.drop_duplicates()`
`length_rating_df.head()`

Out[14]:

| | <code>id</code> | <code>runtime</code> | <code>vote_average</code> | <code>vote_count</code> |
|---|-----------------|----------------------|---------------------------|-------------------------|
| 0 | 27 | NaN | 4.9 | 170 |
| 1 | 90 | 96.0 | 7.1 | 1827 |
| 2 | 93 | 110.0 | 7.9 | 359 |
| 3 | 95 | 116.0 | 6.7 | 4267 |
| 4 | 189 | 94.0 | 6.3 | 2210 |

In [15]: ➔ `# check for null values`
`print(length_rating_df.isnull().sum())`

| | |
|---------------------------|---|
| <code>id</code> | 0 |
| <code>runtime</code> | 1 |
| <code>vote_average</code> | 0 |
| <code>vote_count</code> | 0 |
| <code>dtype: int64</code> | |

In [16]: ➔ `# drop null values`
`length_rating_df.dropna(subset=["runtime"], inplace=True)`
`print(length_rating_df["runtime"].isnull().sum())`

0

In [17]: ➔ `# filter votes`
`vote_limit = 3`
`filtered_movies_df = length_rating_df[length_rating_df['vote_count']`

In [18]: ➔ `# categorize the length of movies`
`movie_length_limits = [0,90,110,140,500]`
`category_names = ["< 90 min (Short)", "90-110 min (Average)", "110-130 min (Long)"]`
`filtered_movies_df["runtime_category"] = pd.cut(`
 `filtered_movies_df["runtime"],`
 `bins=movie_length_limits,`
 `labels=category_names,`
 `right=False`
`)`

```
In [19]: # group runtime category and vote_average
ratings_by_length = filtered_movies_df.groupby("runtime_category", ob
ratings_by_length.head()
```

Out[19]:

| | runtime_category | vote_average |
|---|----------------------|--------------|
| 0 | > 130 min (Epic) | 7.766667 |
| 1 | < 90 min (Short) | 7.625000 |
| 2 | 90-110 min (Average) | 7.466667 |
| 3 | 110-130 min (Long) | 7.400000 |

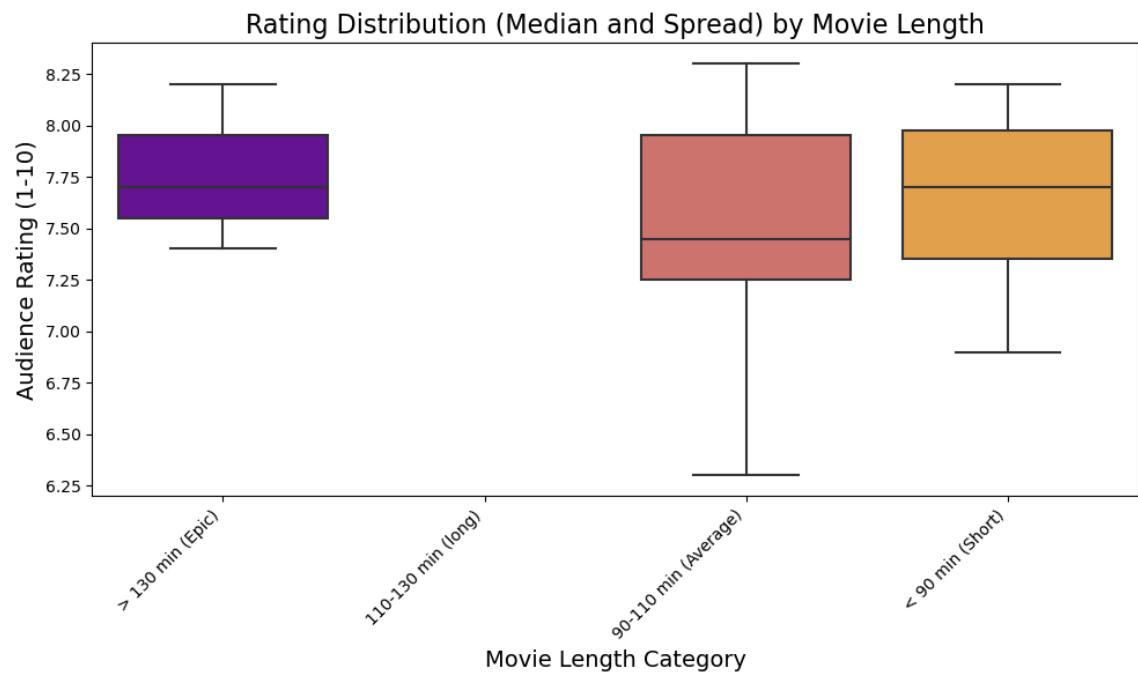
SECTION 3: MOVIE RATINGS BY LENGTH OF MOVIES

This analysis gives a comparison of **movie ratings and the length of the movies**. It shows how the variables runtime and vote_average are related in the dataset. A boxplot was used to show the distribution of votes in relation to runtime. A line graph was also created for the same data

```
In [20]: # use boxplot to see distribution of runtime compared to vote_average
plt.figure(figsize=(10, 6))

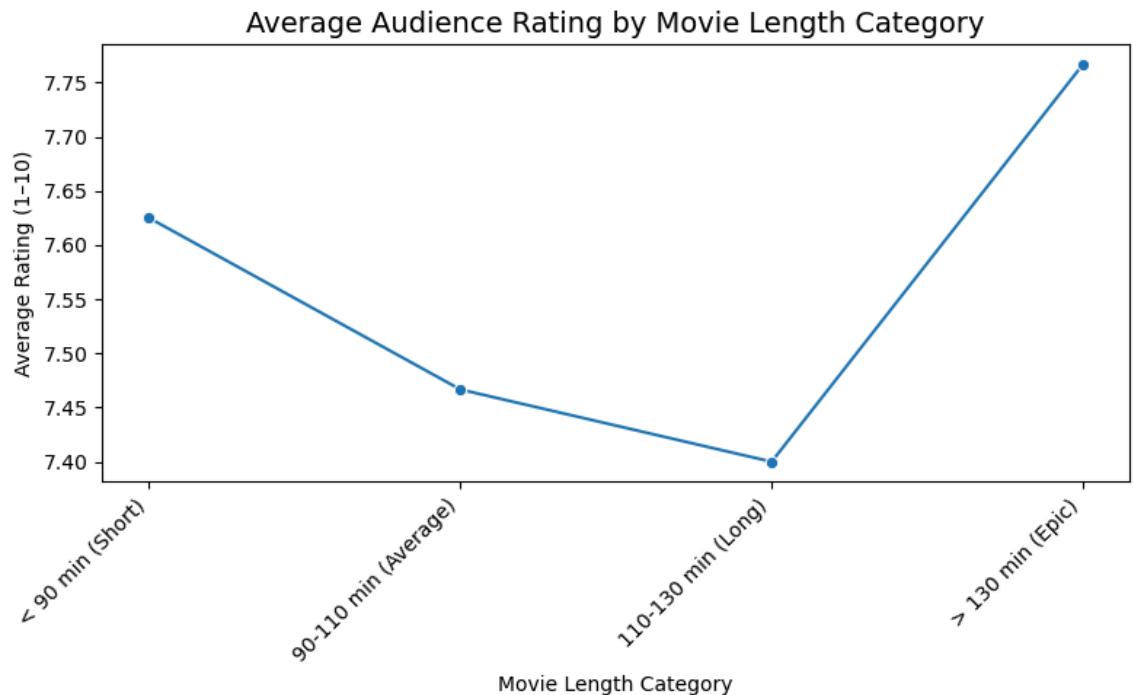
sns.boxplot(
    x="runtime_category",
    y="vote_average",
    data=filtered_movies_df,
    palette="plasma",
    order=['> 130 min (Epic)', '110-130 min (long)', '90-110 min (Average)', '< 90 min (short)'])

plt.title("Rating Distribution (Median and Spread) by Movie Length",
          fontsize=14)
plt.xlabel("Movie Length Category", fontsize=14)
plt.ylabel("Audience Rating (1-10)", fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



```
In [21]: # create plot to check
# average ratings per runtime category
ratings_by_length = (
    filtered_movies_df
    .groupby("runtime_category", observed=True)[ "vote_average" ]
    .mean()
    .reset_index()
)

plt.figure(figsize=(8,5))
sns.lineplot(
    data=ratings_by_length,
    x="runtime_category",
    y="vote_average",
    marker="o"
)
plt.title("Average Audience Rating by Movie Length Category", fontsize=14)
plt.xlabel("Movie Length Category")
plt.ylabel("Average Rating (1-10)")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()
plt.show()
```



SECTION 4: SUCCESSFULL DIRECTORS AND WRITERS BASED ON MOVIE PROFITS AND RATINGS

This part is a data preparation and an analysis to identify **directors and writers** who make successful movies based on the profits they bring to their studios and also based on the `vote_counts` of their movies. It aims to give the new head of the movie studio an idea of the kind of directors and writers to hire.

In [22]: ► # Calculate Total Profit per Director
 director_total_profit = merged_movie_info.groupby('director')['profit']
 # Calculate Average Profit per Director
 director_avg_profit = merged_movie_info.groupby('director')['profit']
 director_avg_profit

Out[22]:

| | director | profit |
|----|---------------------|--------------|
| 0 | Gary Wheeler | 1.684975e+08 |
| 1 | Clint Eastwood | 1.408124e+08 |
| 2 | Bryan Singer | 9.677964e+07 |
| 3 | George Hickenlooper | 9.677964e+07 |
| 4 | Jim Jarmusch | 9.677964e+07 |
| 5 | Joan Chen | 9.677964e+07 |
| 6 | Peter Webber | 9.677964e+07 |
| 7 | Sam Mendes | 9.677964e+07 |
| 8 | Sylvester Stallone | 9.677964e+07 |
| 9 | John Krasinski | 1.552795e+07 |
| 10 | Trey Edward Shults | 1.223359e+07 |

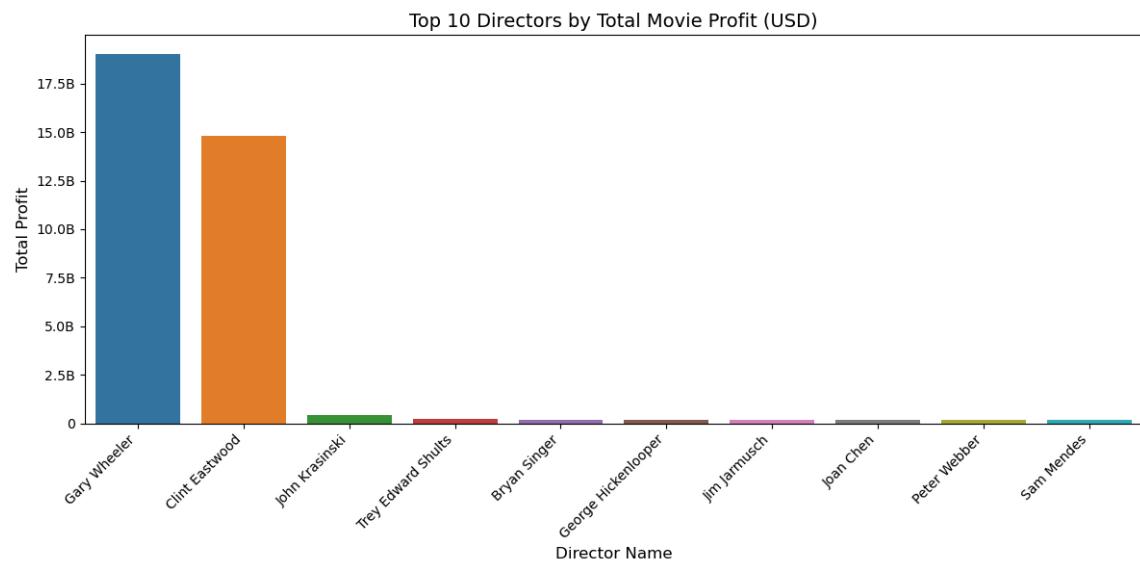
In [23]: ► # Functions to change the format for the plots so as to be able to use
 def format_billion(x, pos):
 """Formatter function to display y-axis ticks in billions (B) or
 if x >= 1e9:
 return f'{x / 1e9:.1f}B'
 elif x >= 1e6:
 return f'{x / 1e6:.0f}M'
 return f'{x:.0f}'

 def format_million(x, pos):
 """Formatter function to display y-axis ticks in millions (M) or
 if x >= 1e9:
 return f'{x / 1e9:.1f}B'
 elif x >= 1e6:
 return f'{x / 1e6:.0f}M'
 return f'{x:.0f}'

 from matplotlib.ticker import FuncFormatter

 # Set up formatter
 billion_formatter = FuncFormatter(format_billion)
 million_formatter = FuncFormatter(format_million)

```
In [24]: # Top 10 Directors by Total Profit ---  
top_10_directors = director_total_profit.head(10)  
  
plt.figure(figsize=(12, 6))  
ax = sns.barplot(x='director', y='profit', data=top_10_directors)  
  
# Apply formatter for billions/millions  
ax.yaxis.set_major_formatter(billion_formatter)  
  
# Labels and formatting  
plt.title('Top 10 Directors by Total Movie Profit (USD)', fontsize=14)  
plt.xlabel('Director Name', fontsize=12)  
plt.ylabel('Total Profit', fontsize=12)  
plt.xticks(rotation=45, ha='right') # Rotate names to prevent overlap  
plt.tight_layout()  
plt.show()
```

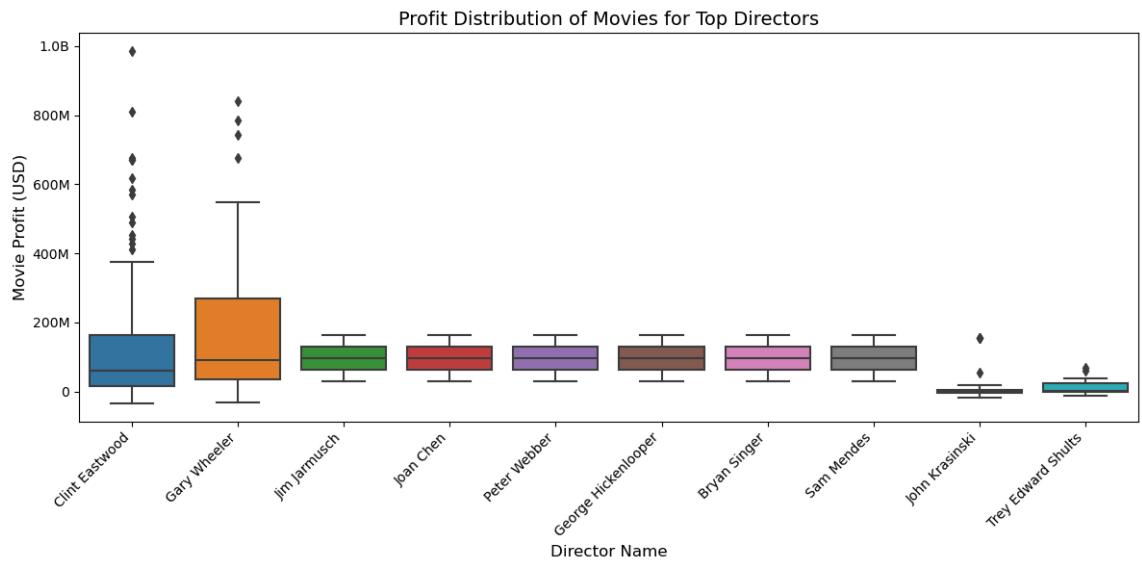


```
In [25]: # Filter data for Top 10 Directors
top_directors_list = top_10_directors['director'].tolist()
filtered_df = merged_movie_info[merged_movie_info['director'].isin(to

# Create boxplot (shows distribution/risk of profits per director)
plt.figure(figsize=(12, 6))
ax = sns.boxplot(x='director', y='profit', data=filtered_df)

# Apply formatter for billions/millions
ax.yaxis.set_major_formatter(billion_formatter)

# Labels and formatting
plt.title('Profit Distribution of Movies for Top Directors', fontsize=14)
plt.xlabel('Director Name', fontsize=12)
plt.ylabel('Movie Profit (USD)', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



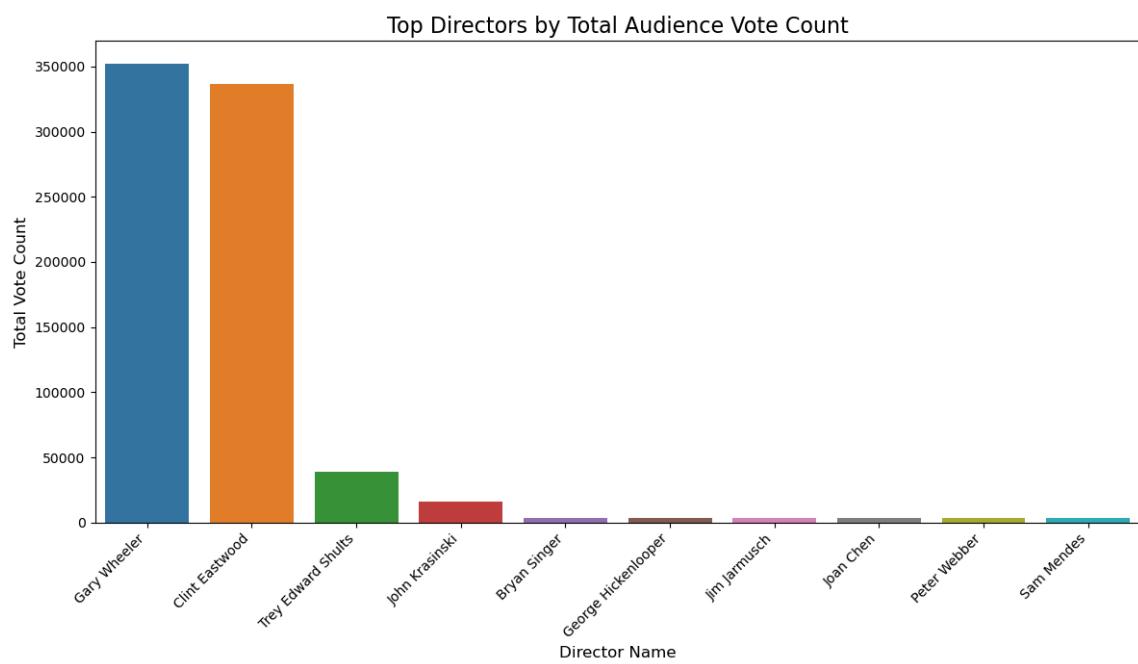
```
In [26]: # Group by director and sum the total vote count and sort
director_total_votes = merged_movie_info.groupby('director')['vote_count'].sum()
director_total_votes.columns = ['Director Name', 'Total Vote Count']

# Select the top directors
top_directors_votes = director_total_votes.head(10)

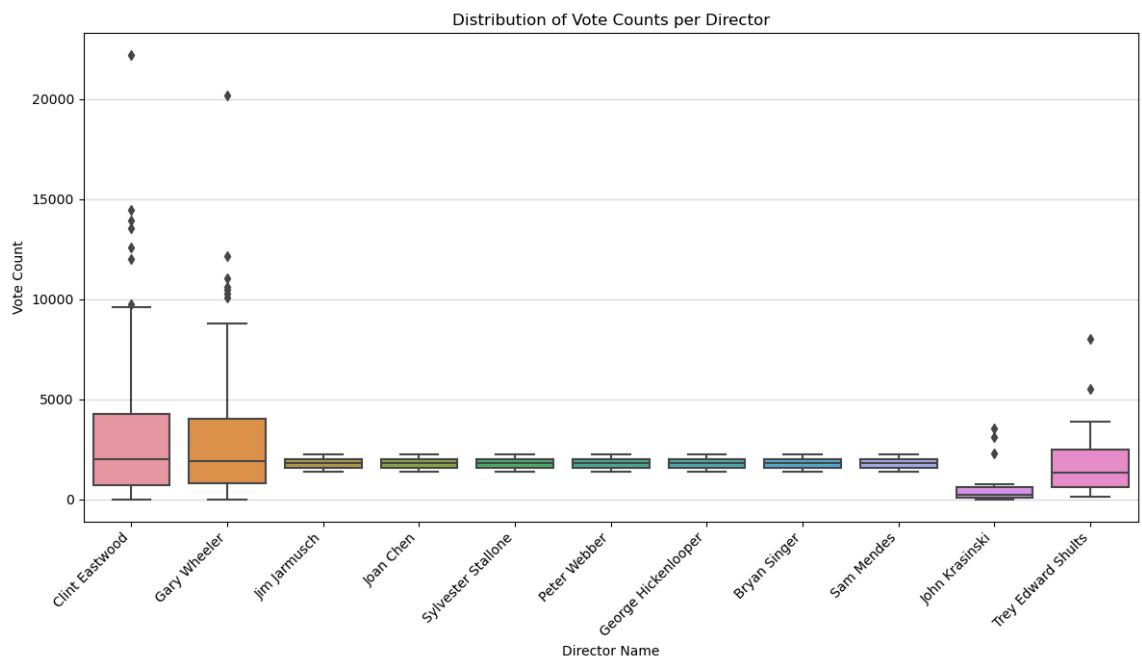
# Bar Chart for Top Directors by Total Vote Count
plt.figure(figsize=(12, 7))

sns.barplot(x='Director Name', y='Total Vote Count', data=top_directors_votes)

plt.title('Top Directors by Total Audience Vote Count', fontsize=16)
plt.xlabel('Director Name', fontsize=12)
plt.ylabel('Total Vote Count', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
```



```
In [27]: # Boxplot on distribution of vote counts per director  
  
plt.figure(figsize=(12, 7))  
  
sns.boxplot(x='director', y='vote_count', data=merged_movie_info)  
  
plt.title('Distribution of Vote Counts per Director')  
plt.xlabel('Director Name')  
plt.ylabel('Vote Count')  
plt.xticks(rotation=45, ha='right')  
plt.grid(axis='y', alpha=0.5)  
plt.tight_layout()
```



```
In [28]: # Distribution of Top Writers by Total Movie Profits
# Split the combined writer names into individual rows
writers_split = merged_movie_info.assign(
    writer=merged_movie_info['writer'].str.split('|'))
.explode('writer')

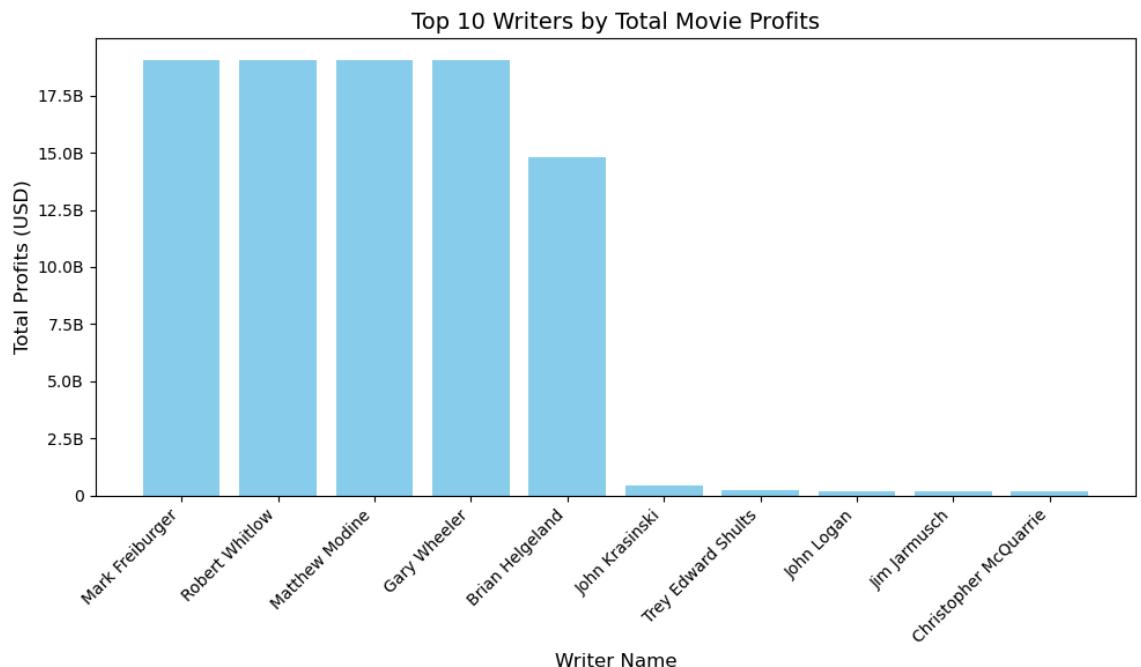
# Clean up whitespace or missing entries
writers_split['writer'] = writers_split['writer'].str.strip()
writers_split = writers_split.dropna(subset=['writer'])

writer_profit = (
    writers_split.groupby('writer')['profit']
    .sum()
    .sort_values(ascending=False)
    .head(10)
    .reset_index()
)

# Create bar plot
plt.figure(figsize=(10,6))
ax = plt.bar(writer_profit['writer'], writer_profit['profit'], color='lightblue')

# Apply billions/millions formatter
plt.gca().yaxis.set_major_formatter(billion_formatter)

# Labels and formatting
plt.title('Top 10 Writers by Total Movie Profits', fontsize=14)
plt.xlabel('Writer Name', fontsize=12)
plt.ylabel('Total Profits (USD)', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



In [29]: # Check Top 10 Writers

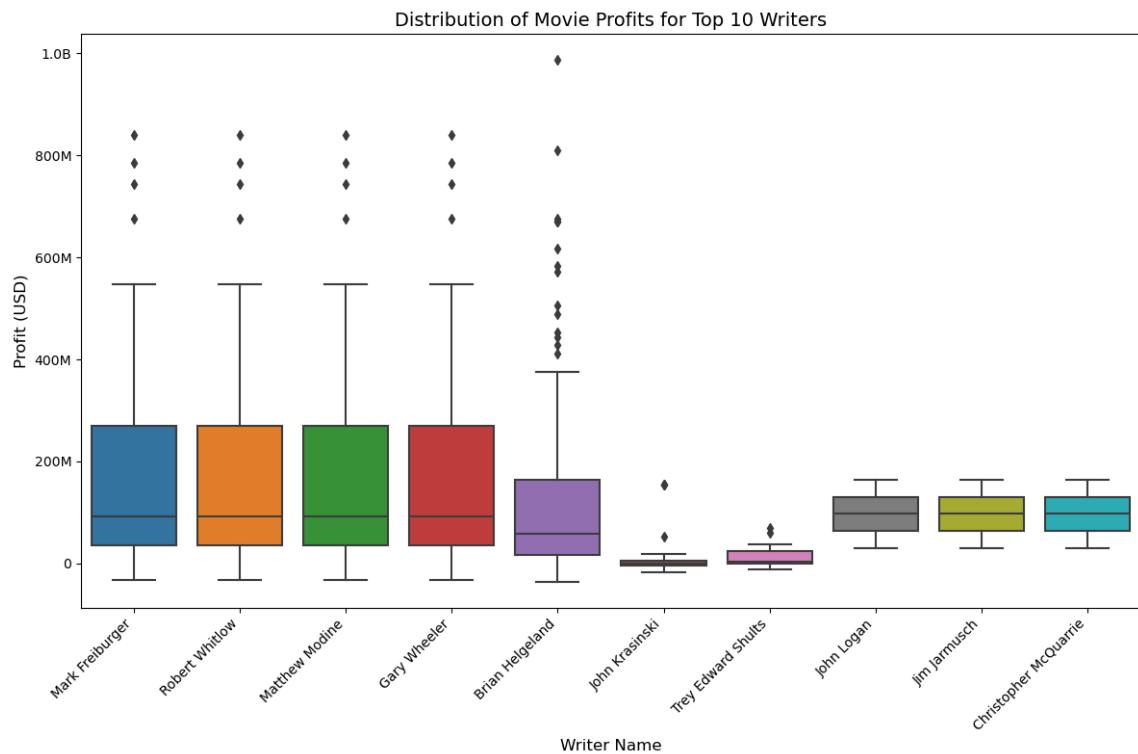
```
# Check Top 10 Writers
top_10_writers = (
    writers_split.groupby('writer')['profit']
    .sum()
    .sort_values(ascending=False)
    .head(10)
    .index.tolist()
)

# Include only movies from the Top 10 Writers
top_writers_data = writers_split[writers_split['writer'].isin(top_10_writers)]

# Create the boxplot
plt.figure(figsize=(12, 8))
ax = sns.boxplot(x='writer', y='profit', data=top_writers_data, order=top_10_writers)

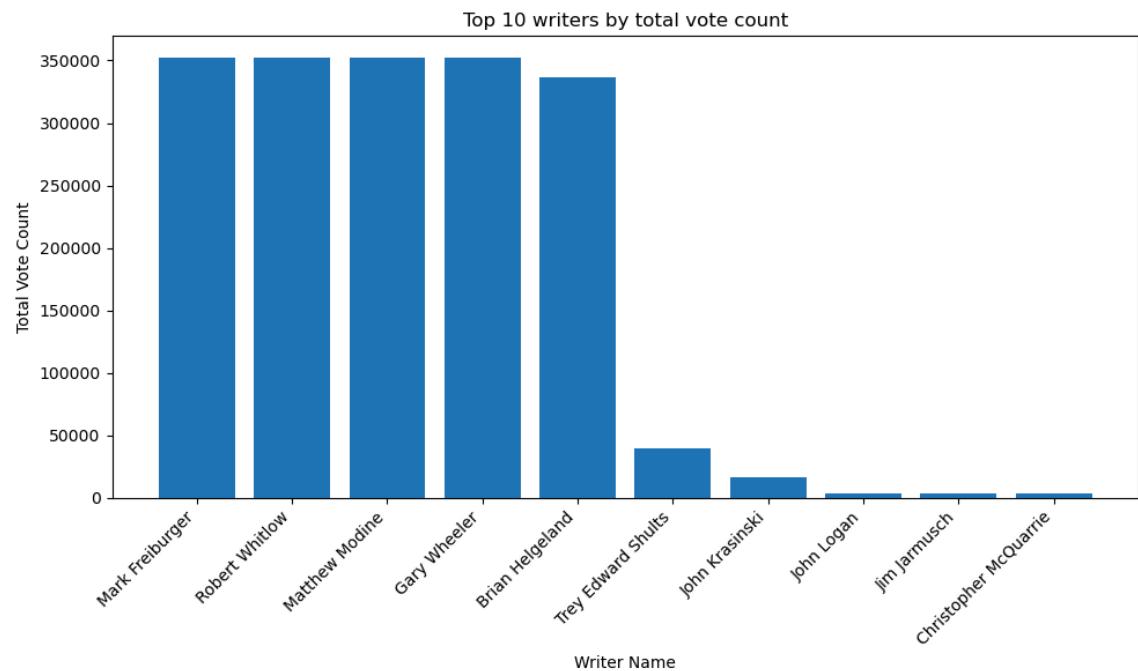
# Apply billions/millions formatter
ax.yaxis.set_major_formatter(billion_formatter)

# Labels and formatting
plt.title('Distribution of Movie Profits for Top 10 Writers', fontsize=14)
plt.xlabel('Writer Name', fontsize=12)
plt.ylabel('Profit (USD)', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



```
In [30]: # writer vs vote counts
writer_votes = (
    writers_split.groupby('writer')['vote_count']
    .sum()
    .sort_values(ascending=False)
    .head(10)
    .reset_index()

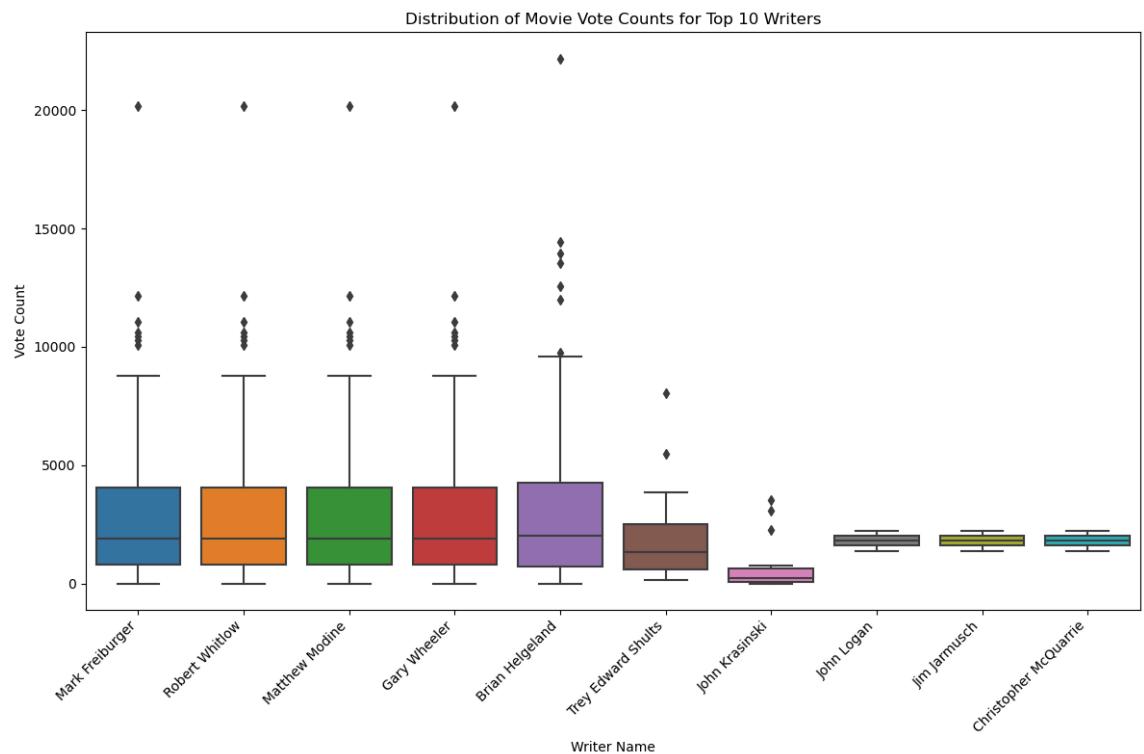
)
# plot
plt.figure(figsize=(10,6))
plt.bar(writer_votes['writer'],writer_votes['vote_count'])
plt.title('Top 10 writers by total vote count')
plt.xlabel('Writer Name')
plt.ylabel('Total Vote Count')
plt.xticks(rotation=45,ha='right')
plt.tight_layout()
plt.show()
```



```
In [31]: # create notebook
top_10_writers_by_votes = (writers_split.groupby('writer')[['vote_count']].sum()
                           .sort_values(ascending=False)
                           .head(10)
                           .index.tolist()
                           )

# Include only movies from the top 10 writers
# This provides the distribution of individual vote counts for the best 10 writers
top_writers_vote_data = writers_split[writers_split['writer'].isin(top_10_writers_by_votes)]

# Create the boxplot
plt.figure(figsize=(12, 8))
sns.boxplot(x='writer', y='vote_count', data=top_writers_vote_data, color='lightgray')
plt.title('Distribution of Movie Vote Counts for Top 10 Writers')
plt.xlabel('Writer Name')
plt.ylabel('Vote Count')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



SECTION 5: MOVIE GENRE ANALYSIS BY RATINGS AND PROFITS

This section shows data preparation and Analysis to establish which **genres have the highest ratings** & to assess which **genres generate the highest gross revenue and profit** to give our stakeholders an idea on which genres they can focus on when they start creating movies.

In [32]: # check for null values in the dataset
df.isna().sum()

Out[32]:

| movie_titles | 1 |
|----------------------|-------|
| genres | 146 |
| original_language | 157 |
| release_date | 1 |
| studio | 44280 |
| production_budget | 1 |
| domestic_gross | 44279 |
| foreign_gross | 44475 |
| worldwide_gross | 44279 |
| revenue | 1 |
| runtime | 394 |
| popularity | 146 |
| production_companies | 146 |
| production_countries | 146 |
| vote_average | 146 |
| vote_count | 146 |
| dtype: | int64 |

In [33]: # Since revenue is the same as gross we will drop the domestic,
foreign and worldwide gross columns since they have a lot of
missing values and the revenue column has less missing values
Drop all the columns not needed in this analysis.
columns_to_drop = [
 'domestic_gross',
 'foreign_gross',
 'worldwide_gross',
 'studio',
 'original_language',
]
df = df.drop(columns=columns_to_drop)

In [34]: # Convert all values to string, then remove \$ and comma, and finally
def clean_currency_column(series):
 cleaned_series = series.astype(str).str.replace('\$', '', regex=False)
 return pd.to_numeric(cleaned_series, errors='coerce')

Parse JSON-like string in the 'genres' column
def extract_genre_names(genres_str):
 if pd.isna(genres_str) or genres_str in ('[]', ''):
 return []
 try:
 # Use simple string replacement to fix single quotes to double
 genres_str_fixed = genres_str.replace("'", "")
 genres_list = json.loads(genres_str_fixed)
 return [g['name'] for g in genres_list]
 except (TypeError, json.JSONDecodeError):
 return []

```
In [35]: df['revenue'] = clean_currency_column(df['revenue']) # Worldwide revenue
df['production_budget'] = clean_currency_column(df['production_budget'])
df['vote_average'] = pd.to_numeric(df['vote_average'], errors='coerce')

# Calculate Profit
df['profit'] = df['revenue'] - df['production_budget']

# Drop rows where key financial data is missing or zero (not reliable)
df.dropna(subset=['revenue', 'production_budget', 'profit', 'vote_average'])
df = df[(df['revenue'] > 1e4) & (df['production_budget'] > 1e4)] # Filter out low-budget movies

# Prepare Genres (Exploding the DataFrame)
df['Genres'] = df['genres'].apply(extract_genre_names)
# Filter out movies with no valid genre
df_genre = df[df['Genres'].apply(lambda x: len(x) > 0)].copy()
# Explode the DataFrame to have one row per movie-genre combination
df_exploded = df_genre.explode('Genres')
```

```
In [36]: # Season function so as to work with the seasons instead of months
df_exploded['release_date'] = pd.to_datetime(df_exploded['release_date'])
df_exploded.dropna(subset=['release_date'], inplace=True) # drop null values

def get_season(month):
    """Maps month number to a cinematic release season."""
    if month in [12, 1, 2]: # December, January, February
        return 'Winter'
    elif month in [3, 4, 5]: # March, April, May
        return 'Spring'
    elif month in [6, 7, 8]: # June, July, August
        return 'Summer'
    else: # September, October, November
        return 'Fall'

df_exploded['Season'] = df_exploded['release_date'].dt.month.apply(get_season)
```

```
In [37]: sns.set_theme(style="whitegrid") # so as to have a standard theme in

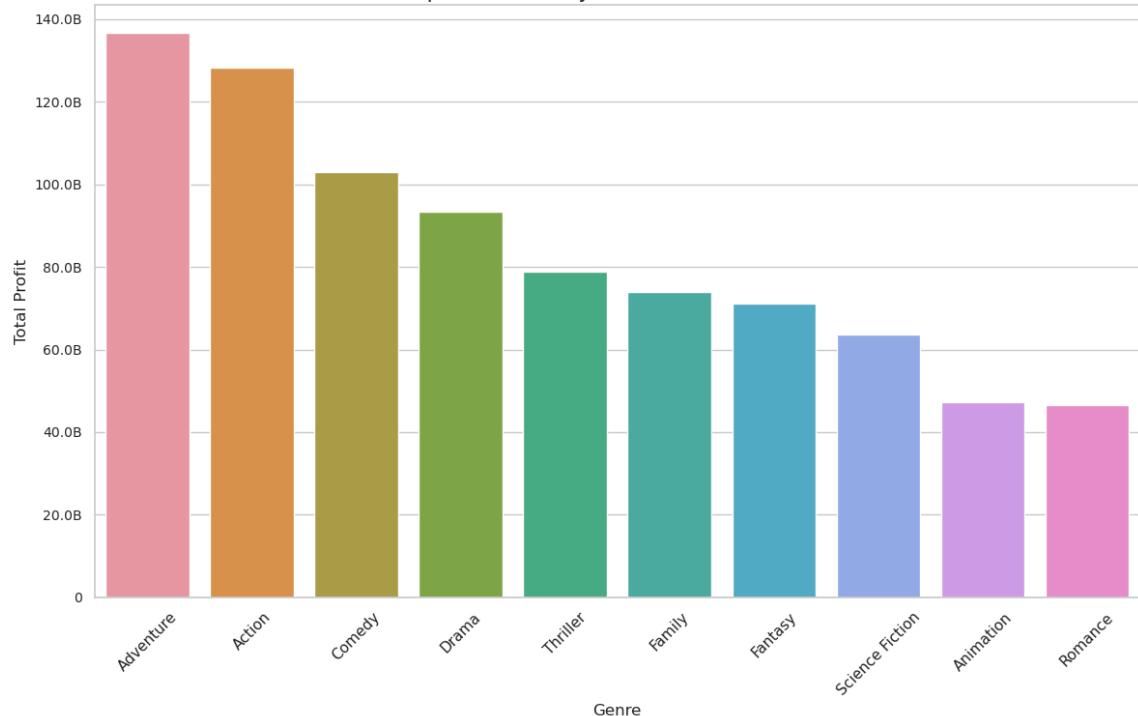
# 1. Genre vs Profit (Total and Average)
# Total Profit
genre_total_profit = df_exploded.groupby('Genres')['profit'].sum().sort_values(ascending=False)

# Average Profit
genre_avg_profit = df_exploded.groupby('Genres')['profit'].mean().sort_values(ascending=False)

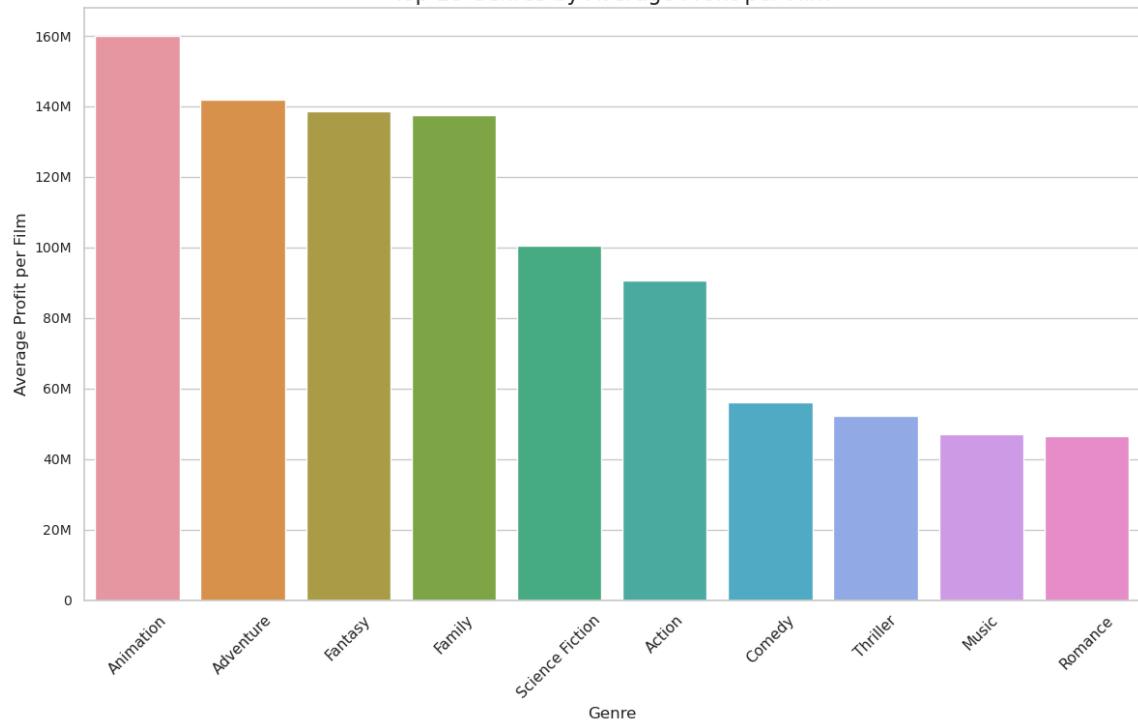
# Plot1. Total Profit
plt.figure(figsize=(12,8))
sns.set_palette("crest")
sns.barplot(x=genre_total_profit.index, y=genre_total_profit.values)
plt.title('Top 10 Genres by Total Worldwide Profit', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Total Profit')
plt.gca().yaxis.set_major_formatter(billion_formatter)
plt.xticks(rotation=45)
plt.yticks(fontsize=10)
plt.tight_layout()
#plt.savefig('images/genre_total_profit.png')
plt.show()

# Plot 2: Average Profit
plt.figure(figsize=(12,8))
sns.set_palette("crest")
sns.barplot(x=genre_avg_profit.index, y=genre_avg_profit.values)
plt.title('Top 10 Genres by Average Profit per Film', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Average Profit per Film')
plt.gca().yaxis.set_major_formatter(million_formatter)
plt.xticks(rotation=45)
plt.yticks(fontsize=10)
plt.tight_layout()
#plt.savefig('images/genre_avg_profit.png')
plt.show()
```

Top 10 Genres by Total Worldwide Profit



Top 10 Genres by Average Profit per Film



In [38]: # 2 Genre vs Revenue (Total and Average)

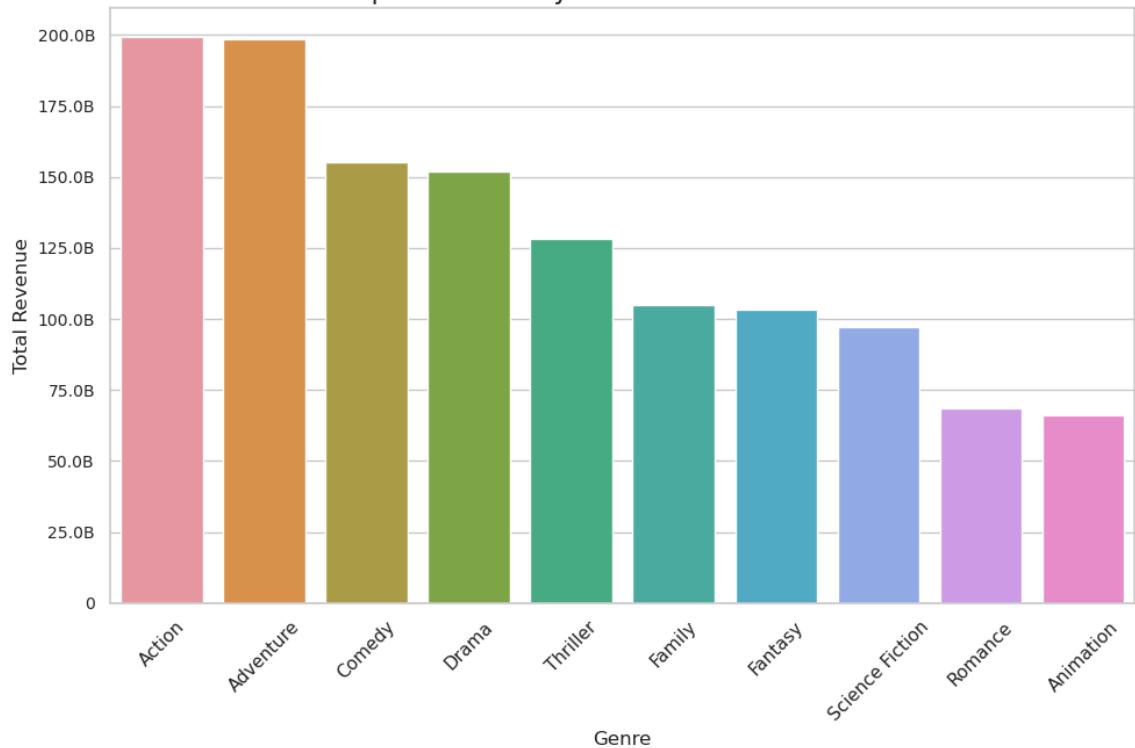
```
# Total Revenue
genre_total_revenue = df_exploded.groupby('Genres')['revenue'].sum()

# Average Revenue
# Re-calculate average on unique movies per genre to avoid skewing the results
# (Although df_exploded.groupby().mean() is statistically okay, this is more accurate)
unique_titles = df_genre.explode('Genres')
genre_avg_revenue = unique_titles.groupby('Genres')['revenue'].mean()

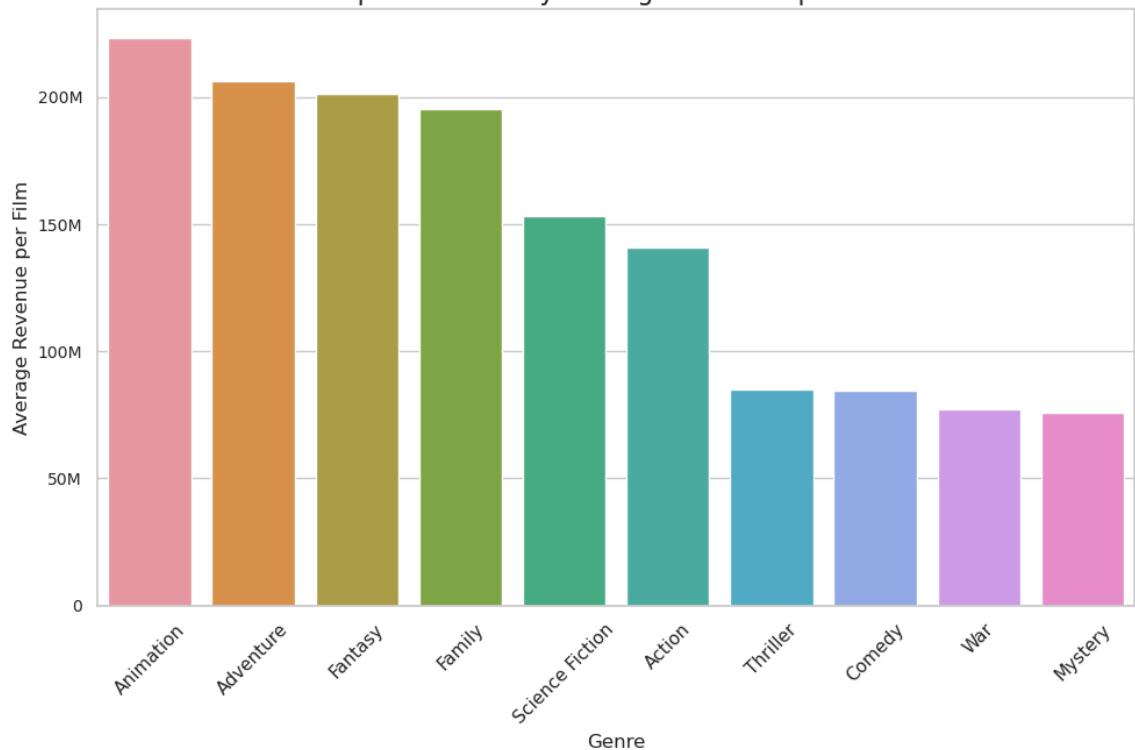
# Plot 1: Total Revenue
plt.figure(figsize=(10,7))
sns.set_palette("crest")
sns.barplot(x=genre_total_revenue.index, y=genre_total_revenue.values)
plt.title('Top 10 Genres by Total Worldwide Revenue', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Total Revenue')
plt.gca().yaxis.set_major_formatter(billion_formatter)
plt.xticks(rotation=45)
plt.yticks(fontsize=10)
plt.tight_layout()
# plt.savefig('images/genre_vs_total_revenue.png')
plt.show()

# Plot 2: Average Revenue
plt.figure(figsize=(10,7))
sns.set_palette("crest")
sns.barplot(x=genre_avg_revenue.index, y=genre_avg_revenue.values)
plt.title('Top 10 Genres by Average Revenue per Film', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Average Revenue per Film')
plt.gca().yaxis.set_major_formatter(million_formatter)
plt.xticks(rotation=45)
plt.yticks(fontsize=10)
plt.tight_layout()
# plt.savefig('images/genre_vs_revenue_avg.png')
plt.show()
```

Top 10 Genres by Total Worldwide Revenue



Top 10 Genres by Average Revenue per Film



In [39]: # 3 Genre vs Ratings (Mean Rating and Movie Count)

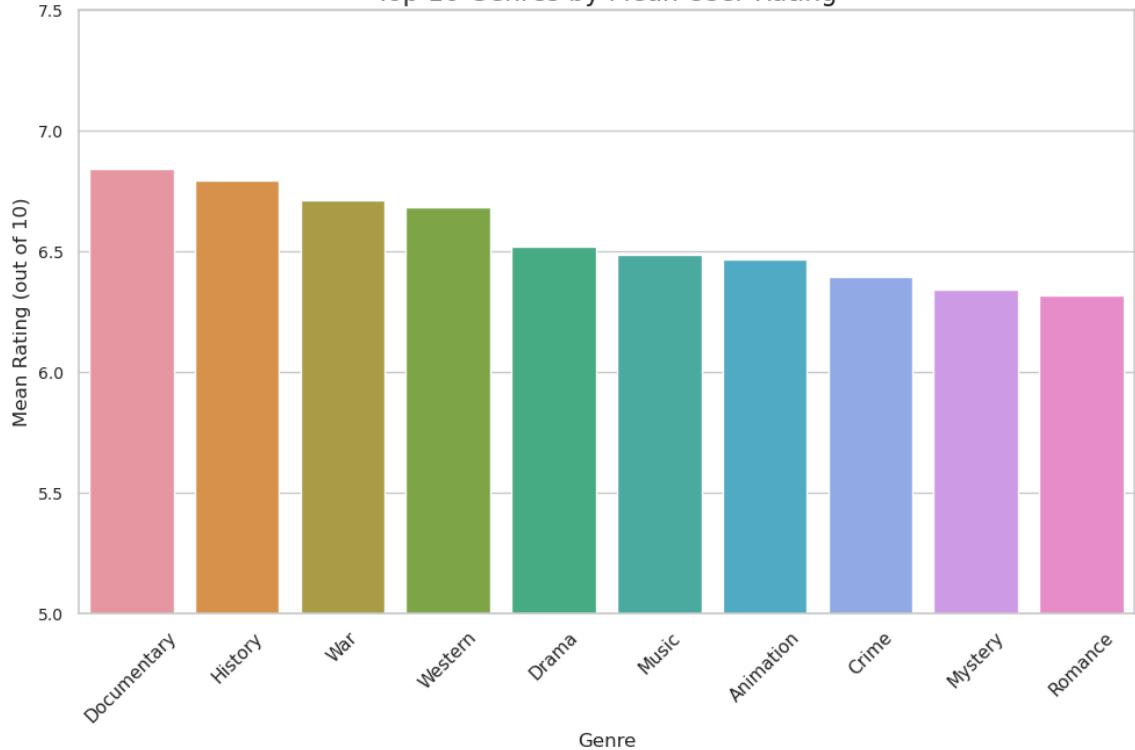
```
# Mean Rating
genre_mean_rating = df_exploded.groupby('Genres')['vote_average'].mean()
genre_mean_rating = genre_mean_rating.sort_values(ascending=False).head(10)

# Movie Count
genre_count = df_exploded.groupby('Genres')['movie_titles'].nunique()

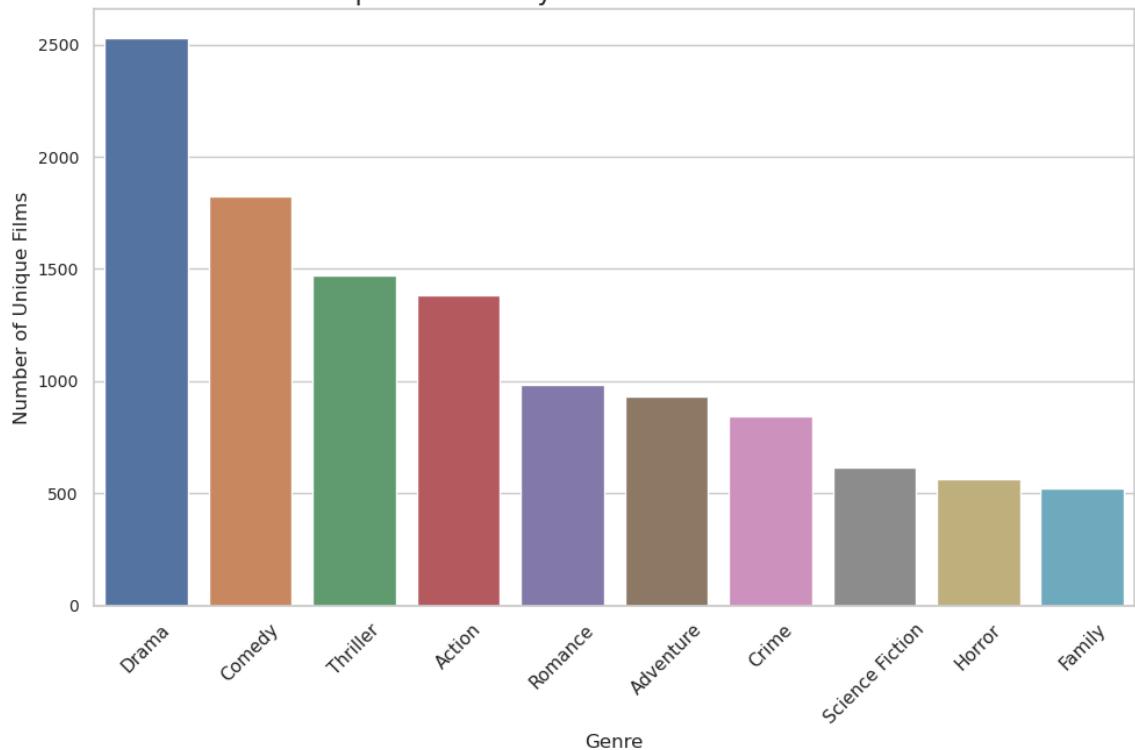
# Plot 1: Mean Rating
plt.figure(figsize=(10, 7))
sns.set_palette("flare")
sns.barplot(x=genre_mean_rating.index, y=genre_mean_rating.values)
plt.title('Top 10 Genres by Mean User Rating', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Mean Rating (out of 10)')
plt.xticks(rotation=45)
plt.ylim(5, 7.5)
plt.yticks(fontsize=10)
plt.tight_layout()
# plt.savefig('images/genre_vs_ratings_count.png')
plt.show()

# Plot 2: Movie Count
plt.figure(figsize=(10, 7))
sns.set_palette("deep")
sns.barplot(x=genre_count.index, y=genre_count.values)
plt.title('Top 10 Genres by Number of Films Produced', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Number of Unique Films')
plt.xticks(rotation=45)
plt.yticks(fontsize=10)
plt.tight_layout()
# plt.savefig('images/genre_vs_movie_count.png')
plt.show()
```

Top 10 Genres by Mean User Rating



Top 10 Genres by Number of Films Produced



SECTION 6: MOVIE PRODUCTION BUDGET VS. WORLDWIDE REVENUE

The scatter plot below illustrates the relationship between a movie's production budget and its worldwide revenue. The plot shows a positive trend: higher production budgets generally lead to higher worldwide revenues. However, the spread of points around the break-even line indicates

that bigger budgets don't always guarantee profits. as some expensive films still underperform.

```
In [40]: # 4 Budget vs Revenue
df_plot = df.copy()

plt.figure(figsize=(10, 8))
sns.set_palette("viridis")
sns.scatterplot(
    x='production_budget',
    y='revenue',
    data=df_plot,
    alpha=0.6,
    s=80,
    linewidth=0
)

# Break-even line
max_value = max(df_plot['production_budget'].max(), df_plot['revenue'])
plt.plot([1, max_value], [1, max_value], linestyle='--', label='Break-even')

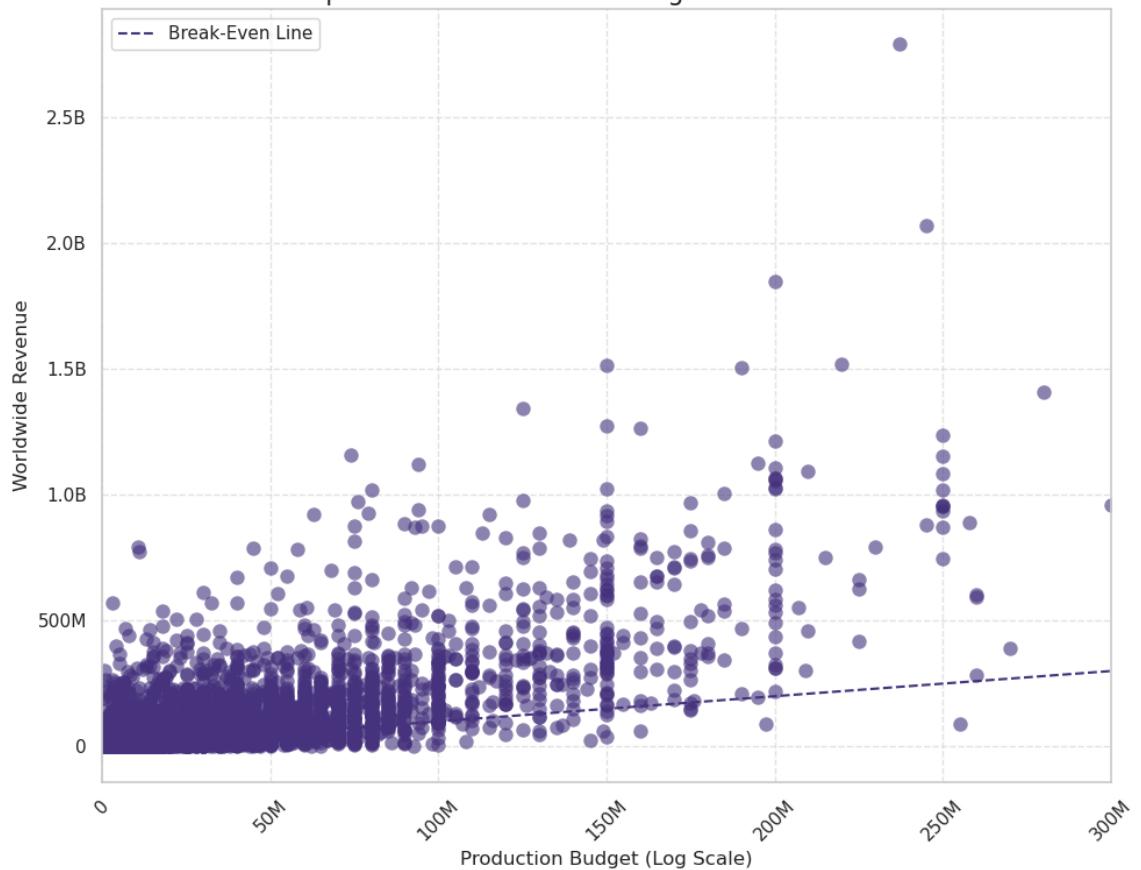
plt.title('Relationship Between Production Budget and Worldwide Revenue')
plt.xlabel('Production Budget (Log Scale)')
plt.ylabel('Worldwide Revenue')

# Apply custom formatters
plt.gca().xaxis.set_major_formatter(billion_formatter)
plt.gca().yaxis.set_major_formatter(billion_formatter)
plt.tick_params(axis='x', rotation=45)

plt.xlim(0, 300_000_000) # Show only budgets ≤ $300M

plt.grid(True, linestyle='--', alpha=0.5)
plt.legend()
plt.tight_layout()
# plt.savefig('images/budget_vs_revenue.png')
plt.show()
```

Relationship Between Production Budget and Worldwide Revenue



SECTION 7: RATINGS VS. PROFITABILITY BY BUDGET CATEGORY

The chart below shows that movies with higher audience ratings generally earn higher profits, especially in the high and blockbuster budget categories. Lower-budget films tend to cluster around smaller profits, while larger-budget films show greater variation indicating that bigger budgets increase both potential reward and financial risk.

```
In [41]: # 5 Ratings vs Profits ( Budget as Hue)

df_plot = df.copy()

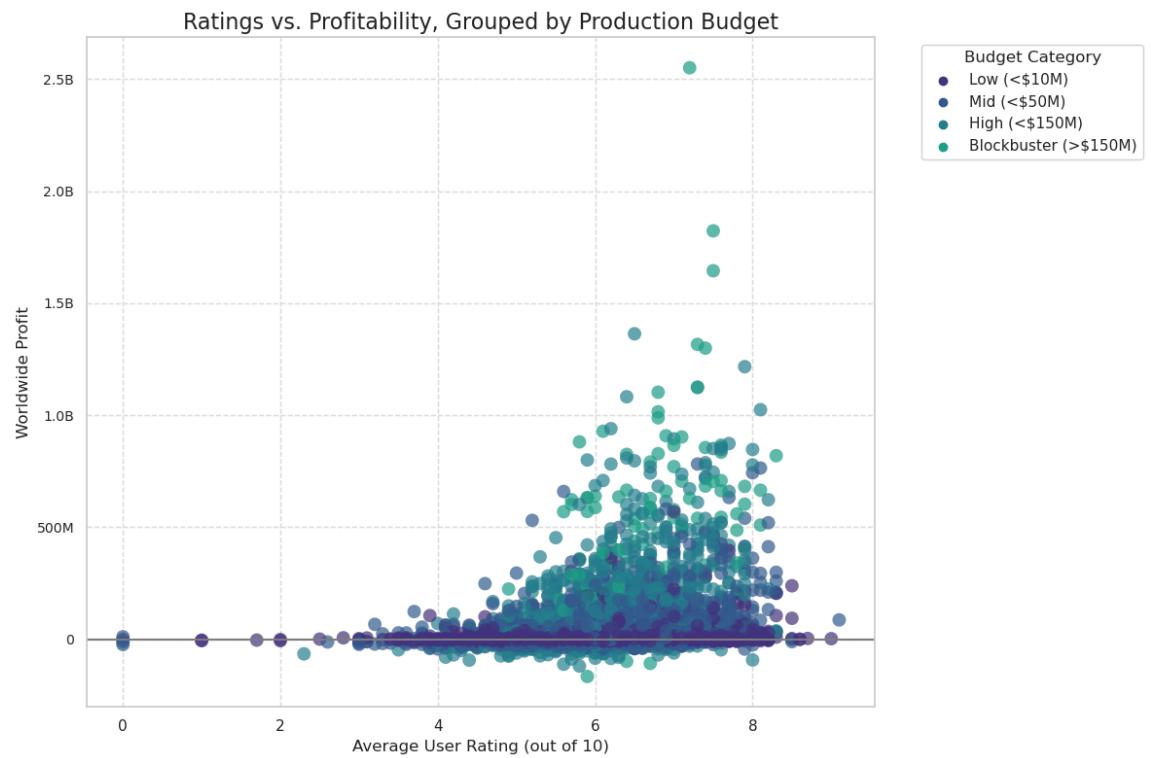
# Create a categorical budget size column for color coding
df_plot['Budget_Category'] = pd.cut(
    df_plot['production_budget'],
    bins=[-np.inf, 1e7, 5e7, 1.5e8, np.inf],
    labels=['Low (<$10M)', 'Mid (<$50M)', 'High (<$150M)', 'Blockbuster']
)
# Scatter plot with Rating on X, Profit on Y, and Budget Category as hue
plt.figure(figsize=(12, 8))
sns.set_palette("viridis")
sns.scatterplot(
    x='vote_average',
    y='profit',
    hue='Budget_Category',
    data=df_plot,
    alpha=0.7,
    s=100,
    linewidth=0,
    hue_order=['Low (<$10M)', 'Mid (<$50M)', 'High (<$150M)', 'Blockbuster']
)

# Draw a line at Profit = 0
plt.axhline(0, color='gray', linestyle='--')

plt.title('Ratings vs. Profitability, Grouped by Production Budget',
          plt.xlabel('Average User Rating (out of 10)')
          plt.ylabel('Worldwide Profit'))

# Apply custom formatter to Y-axis (Profit)
plt.gca().yaxis.set_major_formatter(billion_formatter)
plt.tick_params(axis='y', labelsize=10)

plt.legend(title='Budget Category', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
# plt.savefig('images/ratings_vs_profits_budget.png', dpi=300)
plt.show()
```



SECTION 8: PROFITABILITY BY MOVIE RELEASE SEASON

Movies released in summer and spring tend to achieve the highest profits, likely due to school holidays and major blockbuster releases. In contrast, fall and winter films show lower average profits, suggesting that release timing plays a key role in maximizing box office returns.

In [42]: # 6 Seasons vs Profits

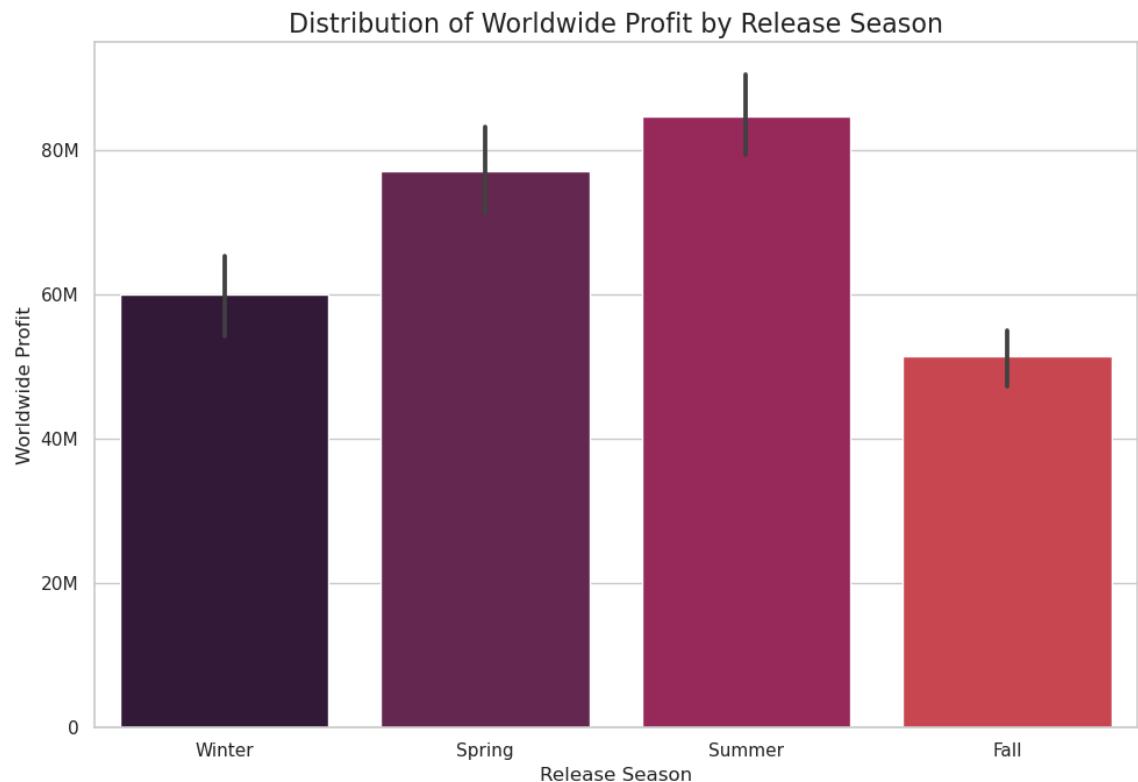
```
season_order = ['Winter', 'Spring', 'Summer', 'Fall']

plt.figure(figsize=(10, 7))
sns.set_palette("rocket")
sns.barplot(
    x='Season',
    y='profit',
    data=df_exploded,
    order=season_order,
)

plt.title('Distribution of Worldwide Profit by Release Season', fontweight='bold')
plt.xlabel('Release Season')
plt.ylabel('Worldwide Profit')

plt.gca().yaxis.set_major_formatter(billion_formatter)

plt.tight_layout()
# plt.savefig('images/seasons_vs_profits.png')
plt.show()
```



SECTION 9: CONCLUSIONS AND RECOMMENDATIONS

CONCLUSIONS

1. **Movie length vs audience ratings:** The longer movies tend to be rated more favorably

2. **Directors and writers Vs Movie Performance:** The directors generating high revenue to be considered for better returns.
3. **Genres effects on ratings, gross revenue and profit:** Genres with high total profit are better options for studios aiming for consistent returns
4. **Movie release season vs profit :** Studios to prioritize Summer and Spring for major release to maximize revenue

RECOMMENDATIONS

1. **Movie length vs audience ratings:** Content creators to consider longer movie run times, they probably give an allowance for deeper story telling.
2. **Directors and writers Vs Movie Performance:** The studio to consider prioritizing directors who have released movies that generated high revenues e.g Clint Eastwood and Gary Wheeler
3. **Genres effects on ratings, gross revenue and profit:** The studios to consider prioritizing genres with high total profit for consistent returns
4. **Movie release season vs profit :** Seasonality should be considered in release planning and marketing strategies.