

movies

November 7, 2025

0.1 BUSINESS UNDERSTANDING

This project analyzes movie performance data to uncover the key factors that drive box office success and audience ratings. By examining variables such as genre, budget, runtime, ratings, and contributor roles (directors and writers), it aims to help studios make data-driven decisions in budgeting, hiring, and marketing. The goal is to identify high-impact talent, understand how production choices affect profitability, and provide actionable insights to improve investment, content, and strategic planning in the film industry.

0.2 SECTION 1: DATA UNDERSTANDING

The datasets used in this project are sourced from publicly available movie databases, including Box Office Mojo (BOM), The Numbers (TN), IMDB database and The Movie Database (TMDb). They provide structured information on thousands of films released over the past two decades. The data spans multiple formats; categorical variables such as genre, director, and writer, numerical variables such as budget, gross revenue, and ratings and temporal variables such as release year and runtime. The dataset includes:

- Film titles & release dates
- Genres (one or multiple per film)
- Production budget & box office revenue
- Ratings & popularity metrics
- Contributor data (directors, writers etc.)
- Movie runtime

```
[1]: # imports
import pandas as pd # cleaning data
import matplotlib.pyplot as plt # data visualization
import seaborn as sns # data visualization
import numpy as np # python calculations
import json # parse data
import warnings # ignore warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
[2]: # import all required datasets
length_data = pd.read_csv("Data/rt.movie_info.tsv", sep='\t')
ratings_data = pd.read_csv("Data/tmdb.movies.csv")
merged_movie_info = pd.read_csv("Data/merged_movie_info.csv", index_col=[0])
df = pd.read_csv('Data/clean_full_movie_data.csv')
```

```
[3]: # check the first rows of dataset
length_data.head()
```

```
[3]: id synopsis rating \
0 1 This gritty, fast-paced, and innovative police... R
1 3 New York City, not-too-distant-future: Eric Pa... R
2 5 Illeana Douglas delivers a superb performance ... R
3 6 Michael Douglas runs afoul of a treacherous su... R
4 7 NaN NR

genre director \
0 Action and Adventure|Classics|Drama William Friedkin
1 Drama|Science Fiction and Fantasy David Cronenberg
2 Drama|Musical and Performing Arts Allison Anders
3 Drama|Mystery and Suspense Barry Levinson
4 Drama|Romance Rodney Bennett

writer theater_date dvd_date currency \
0 Ernest Tidyman Oct 9, 1971 Sep 25, 2001 NaN
1 David Cronenberg|Don DeLillo Aug 17, 2012 Jan 1, 2013 $
2 Allison Anders Sep 13, 1996 Apr 18, 2000 NaN
3 Paul Attanasio|Michael Crichton Dec 9, 1994 Aug 27, 1997 NaN
4 Giles Cooper NaN NaN NaN

box_office runtime studio
0 NaN 104 minutes NaN
1 600,000 108 minutes Entertainment One
2 NaN 116 minutes NaN
3 NaN 128 minutes NaN
4 NaN 200 minutes NaN
```

```
[4]: # check first rows of dataset
ratings_data.head()
```

```
Unnamed: 0 genre_ids id original_language \
0 0 [12, 14, 10751] 12444 en
1 1 [14, 12, 16, 10751] 10191 en
2 2 [12, 28, 878] 10138 en
3 3 [16, 35, 10751] 862 en
4 4 [28, 878, 12] 27205 en

original_title popularity release_date \
0 Harry Potter and the Deathly Hallows: Part 1 33.533 2010-11-19
1 How to Train Your Dragon 28.734 2010-03-26
2 Iron Man 2 28.515 2010-05-07
3 Toy Story 28.005 1995-11-22
4 Inception 27.920 2010-07-16

title vote_average vote_count
0 Harry Potter and the Deathly Hallows: Part 1 7.7 10788
```

1	How to Train Your Dragon	7.7	7610
2	Iron Man 2	6.8	12368
3	Toy Story	7.9	10174
4	Inception	8.3	22186

[5]: # check dataset information
length_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1560 entries, 0 to 1559
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               1560 non-null    int64  
 1   synopsis         1498 non-null    object  
 2   rating            1557 non-null    object  
 3   genre             1552 non-null    object  
 4   director          1361 non-null    object  
 5   writer            1111 non-null    object  
 6   theater_date     1201 non-null    object  
 7   dvd_date          1201 non-null    object  
 8   currency          340 non-null     object  
 9   box_office        340 non-null     object  
 10  runtime           1530 non-null    object  
 11  studio            494 non-null     object  
dtypes: int64(1), object(11)
memory usage: 146.4+ KB
```

[6]: # check dataset information
ratings_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26517 entries, 0 to 26516
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        26517 non-null    int64  
 1   genre_ids         26517 non-null    object  
 2   id                26517 non-null    int64  
 3   original_language 26517 non-null    object  
 4   original_title     26517 non-null    object  
 5   popularity         26517 non-null    float64 
 6   release_date       26517 non-null    object  
 7   title              26517 non-null    object  
 8   vote_average       26517 non-null    float64 
 9   vote_count          26517 non-null    int64  
dtypes: float64(2), int64(3), object(5)
memory usage: 2.0+ MB
```

```
[7]: # look at the first five rows
merged_movie_info.head()
```

```
[7]:                                     title studio year \
1                           Inception      WB 2010
8 The Chronicles of Narnia: The Voyage of the Da...    Fox 2010
30                      Gulliver's Travels    Fox 2010
34                      Due Date      WB 2010
36                      Yogi Bear      WB 2010

   production_budget worldwide_gross popularity vote_average vote_count \
1           160000000  835524642     27.920       8.3        22186
8           155000000  418186950     17.382       6.3        3196
30          112000000  232017848     10.768       5.1        1282
34          65000000  211739043     12.445       6.3        2973
36          80000000  204774690     9.096       5.3        387

   profit genre director \
1 675524642 Drama|Mystery and Suspense Clint Eastwood
8 263186950 Drama|Mystery and Suspense Gary Wheeler
30 120017848 Drama|Mystery and Suspense Gary Wheeler
34 146739043 Drama|Mystery and Suspense Clint Eastwood
36 124774690 Drama|Mystery and Suspense Clint Eastwood

   writer
1 Brian Helgeland
8 Mark Freiburger|Gary Wheeler|Robert Whitlow|Ma...
30 Mark Freiburger|Gary Wheeler|Robert Whitlow|Ma...
34 Brian Helgeland
36 Brian Helgeland
```

```
[8]: # checkout dataset information
merged_movie_info.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 282 entries, 1 to 2470
Data columns (total 12 columns):
 #  Column            Non-Null Count Dtype
 --- 
 0  title             282 non-null   object
 1  studio            282 non-null   object
 2  year              282 non-null   int64
 3  production_budget 282 non-null   int64
 4  worldwide_gross   282 non-null   int64
 5  popularity         282 non-null   float64
 6  vote_average       282 non-null   float64
 7  vote_count         282 non-null   int64
 8  profit             282 non-null   int64
```

```
9   genre           282 non-null    object
10  director        282 non-null    object
11  writer          282 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 28.6+ KB
```

```
[9]: # check the first rows
df.head()
```

```
[9]: movie_titles \
0          toy story 3
1          inception
2          shrek forever after
3  the twilight saga: eclipse
4          iron man 2

genres original_language \
0  [{}'id': 16, 'name': 'Animation'], {}'id': 10751...      en
1  [{}'id': 28, 'name': 'Action'], {}'id': 53, 'nam...      en
2  [{}'id': 35, 'name': 'Comedy'], {}'id': 12, 'nam...      en
3  [{}'id': 12, 'name': 'Adventure'], {}'id': 14, '...      en
4  [{}'id': 12, 'name': 'Adventure'], {}'id': 28, '...      en

release_date studio production_budget domestic_gross foreign_gross \
0  2010-06-16     BV            2000000000 $415,004,880  652000000
1  2010-07-14     WB            1600000000 $292,576,195  535700000
2  2010-05-16     P/DW          1650000000 $238,736,787  513900000
3  2010-06-23     Sum.          680000000  $300,531,751  398000000
4  2010-04-28     Par.          2000000000 $312,433,331  311500000

worldwide_gross      revenue  runtime popularity \
0  $1,068,879,522  1066969703    103.0  16.966470
1  $835,524,642   825532764     148.0  29.108149
2  $756,244,673   752600867     93.0   11.803808
3  $706,102,828   698491347    124.0   34.047399
4  $621,156,389   623933331    124.0   19.083344

production_companies \
0  [{}'name': 'Walt Disney Pictures', 'id': 2}, {}...
1  [{}'name': 'Legendary Pictures', 'id': 923}, {}...
2  [{}'name': 'DreamWorks Animation', 'id': 521}]
3  [{}'name': 'Summit Entertainment', 'id': 491}, ...
4  [{}'name': 'Marvel Studios', 'id': 420}]

production_countries  vote_average  vote_count
0  [{}'iso_3166_1': 'US', 'name': 'United States o...      7.6      4710.0
1  [{}'iso_3166_1': 'GB', 'name': 'United Kingdom'...      8.1      14075.0
```

```

2  [{"iso_3166_1": "US", "name": "United States o...", 6.0, 2021.0}
3  [{"iso_3166_1": "US", "name": "United States o...", 5.9, 2382.0}
4  [{"iso_3166_1": "US", "name": "United States o...", 6.6, 6969.0}

```

[10]: # check dataset information
df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45749 entries, 0 to 45748
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   movie_titles     45748 non-null   object  
 1   genres           45603 non-null   object  
 2   original_language 45592 non-null   object  
 3   release_date     45748 non-null   object  
 4   studio           1469 non-null   object  
 5   production_budget 45748 non-null   object  
 6   domestic_gross    1470 non-null   object  
 7   foreign_gross     1274 non-null   object  
 8   worldwide_gross   1470 non-null   object  
 9   revenue          45748 non-null   object  
 10  runtime          45355 non-null   float64 
 11  popularity        45603 non-null   float64 
 12  production_companies 45603 non-null   object  
 13  production_countries 45603 non-null   object  
 14  vote_average      45603 non-null   float64 
 15  vote_count         45603 non-null   float64 

dtypes: float64(4), object(12)
memory usage: 5.6+ MB

```

0.3 SECTION 2: DATA PREPARATION

To ensure reliable analysis, the following steps were performed: - Merging of the datasets - Cleaned inconsistent formatting and missing values - Converted monetary values to numeric types - Extracted and separated multiple genres per film - Calculated profit and categorized release dates into seasons

[11]: #merging the ratings and length datasets
length_rating_df = pd.merge(
 length_data[["id", "runtime"]],
 ratings_data[["id", "vote_average", "vote_count"]],
 on="id",
 how="inner"
)
length_rating_df.head()

```
[11]:    id      runtime  vote_average  vote_count
0    27        NaN        4.9         170
1    90  96 minutes        7.1        1827
2    93 110 minutes        7.9         359
3    95 116 minutes        6.7        4267
4   189  94 minutes        6.3        2210
```

```
[12]: # converting the columns to numeric values
length_rating_df["runtime"] = length_rating_df["runtime"].astype(str).str.
    replace(" minutes", "", regex=False).str.strip()
length_rating_df["runtime"] = pd.to_numeric(length_rating_df["runtime"], errors="coerce")
```

```
[13]: # check duplicated values
print(length_rating_df.duplicated().value_counts())
```

```
False    27
True     5
Name: count, dtype: int64
```

```
[14]: #drop duplicates
length_rating_df = length_rating_df.drop_duplicates()
length_rating_df.head()
```

```
[14]:    id      runtime  vote_average  vote_count
0    27        NaN        4.9         170
1    90      96.0        7.1        1827
2    93     110.0        7.9         359
3    95     116.0        6.7        4267
4   189      94.0        6.3        2210
```

```
[15]: # check for null values
print(length_rating_df.isnull().sum())
```

```
id          0
runtime      1
vote_average 0
vote_count    0
dtype: int64
```

```
[16]: # drop null values
length_rating_df.dropna(subset=["runtime"], inplace=True)
print(length_rating_df["runtime"].isnull().sum())
```

```
0
```

```
[17]: # filter votes
vote_limit = 3
```

```
filtered_movies_df = length_rating_df[length_rating_df['vote_count'] ↴>=vote_limit].copy()
```

```
[18]: # categorize the length of movies
movie_length_limits = [0,90,110,140,500]
category_names = ["< 90 min (Short)", "90-110 min (Average)", "110-130 min ↴(Long)", "> 130 min (Epic)"]
filtered_movies_df["runtime_category"] = pd.cut(
    filtered_movies_df["runtime"],
    bins=movie_length_limits,
    labels=category_names,
    right=False
)
```

```
[19]: # group runtime category and vote_average
ratings_by_length = filtered_movies_df.groupby("runtime_category", ↴observed=True)[["vote_average"]].mean().sort_values(ascending=False).reset_index()
ratings_by_length.head()
```

```
[19]:   runtime_category  vote_average
0      > 130 min (Epic)    7.766667
1      < 90 min (Short)    7.625000
2    90-110 min (Average)    7.466667
3    110-130 min (Long)    7.400000
```

0.4 SECTION 3: MOVIE RATINGS BY LENGTH OF MOVIES

This analysis gives a comparison of **movie ratings and the length of the movies**. It shows how the variables runtime and vote_average are related in the dataset. A boxplot was used to show the distribution of votes in relation to runtime. A line graph was also created for the same data

```
[20]: # use boxplot to see distribution of runtime compared to vote_average
plt.figure(figsize=(10, 6))

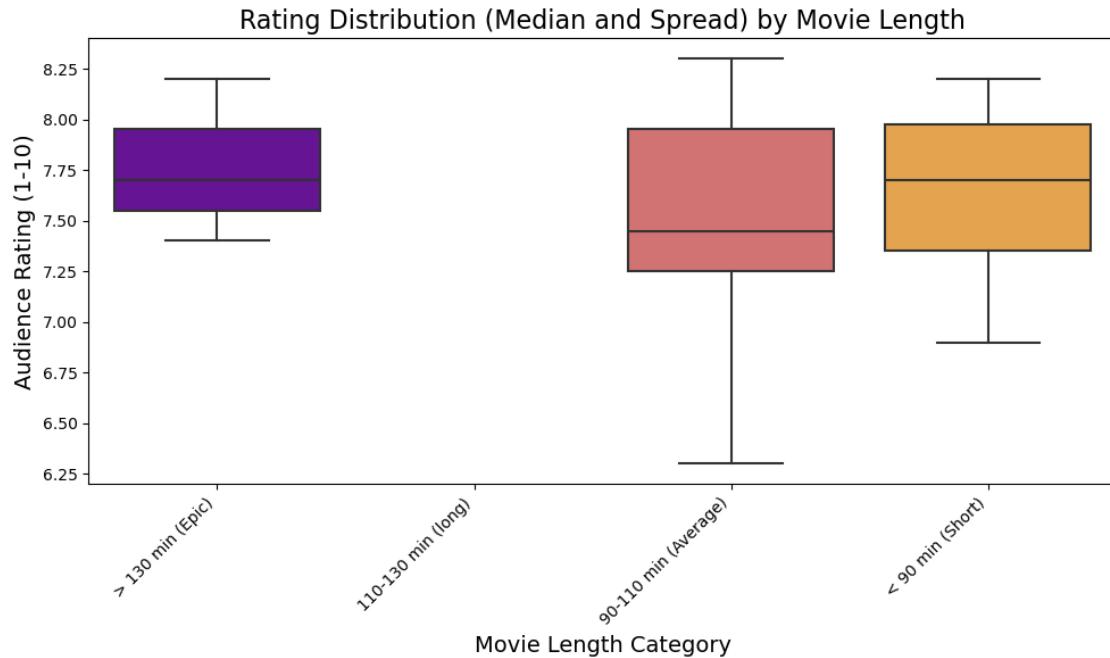
sns.boxplot(
    x="runtime_category",
    y="vote_average",
    data=filtered_movies_df,
    palette="plasma",
    order=['> 130 min (Epic)', '110-130 min (long)', '90-110 min (Average)', '< ↴90 min (Short)']
)

plt.title("Rating Distribution (Median and Spread) by Movie Length", ↴fontsize=16)
plt.xlabel("Movie Length Category", fontsize=14)
```

```

plt.ylabel("Audience Rating (1-10)", fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

```



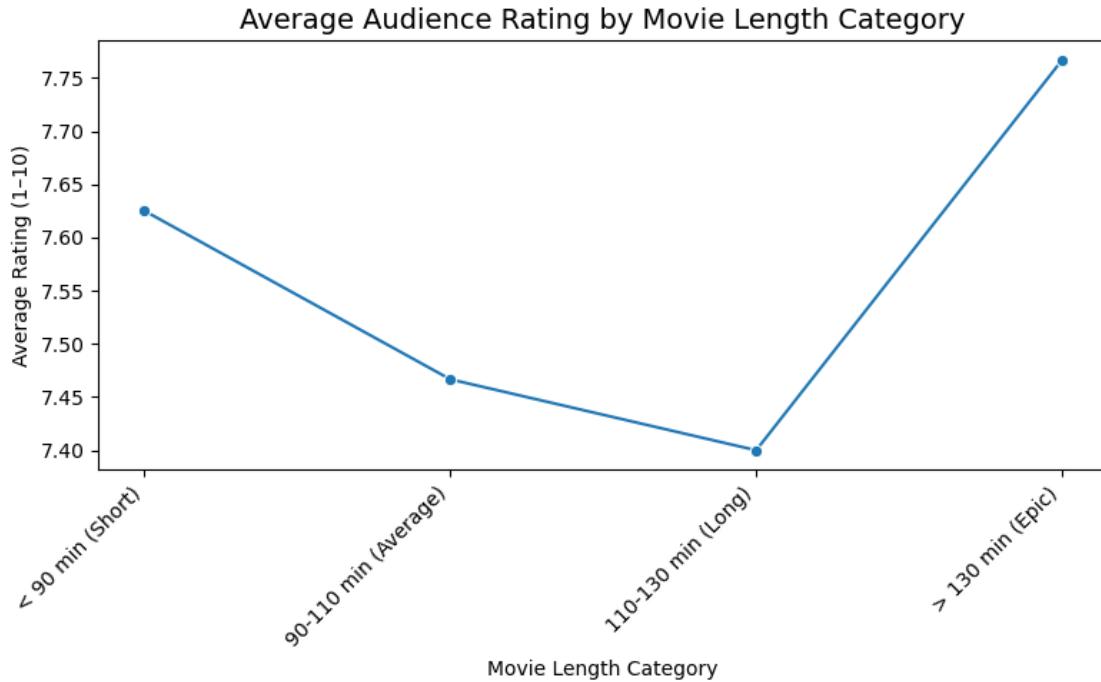
```

[21]: # create plot to check
# average ratings per runtime category
ratings_by_length = (
    filtered_movies_df
    .groupby("runtime_category", observed=True)[ "vote_average"]
    .mean()
    .reset_index()
)

plt.figure(figsize=(8,5))
sns.lineplot(
    data=ratings_by_length,
    x="runtime_category",
    y="vote_average",
    marker="o"
)
plt.title("Average Audience Rating by Movie Length Category", fontsize=14)
plt.xlabel("Movie Length Category")
plt.ylabel("Average Rating (1-10)")
plt.xticks(rotation=45, ha="right")

```

```
plt.tight_layout()
plt.show()
```



0.5 SECTION 4: SUCCESSFULL DIRECTORS AND WRITERS BASED ON MOVIE PROFITS AND RATINGS

This part is a data preparation and an analysis to identify **directors and writers** who make successful movies based on the profits they bring to their studios and also based on the `vote_counts` of their movies. It aims to give the new head of the movie studio an idea of the kind of directors and writers to hire.

```
[22]: # Calculate Total Profit per Director
director_total_profit = merged_movie_info.groupby('director')['profit'].sum().
    sort_values(ascending=False).reset_index()
# Calculate Average Profit per Director
director_avg_profit = merged_movie_info.groupby('director')['profit'].mean().
    sort_values(ascending=False).reset_index()
director_avg_profit
```

```
[22]:      director      profit
0      Gary Wheeler  1.684975e+08
1      Clint Eastwood 1.408124e+08
2      Bryan Singer  9.677964e+07
3  George Hickenlooper  9.677964e+07
4      Jim Jarmusch  9.677964e+07
```

```

5           Joan Chen  9.677964e+07
6           Peter Webber 9.677964e+07
7           Sam Mendes  9.677964e+07
8   Sylvester Stallone 9.677964e+07
9           John Krasinski 1.552795e+07
10          Trey Edward Shults 1.223359e+07

```

```
[23]: # Functions to change the format for the plots so as to be able to us billions
       ↪and millions instead of the math format
def format_billion(x, pos):
    """Formatter function to display y-axis ticks in billions (B) or millions
       ↪(M)."""
    if x >= 1e9:
        return f'{x / 1e9:.1f}B'
    elif x >= 1e6:
        return f'{x / 1e6:.0f}M'
    return f'{x:.0f}'

def format_million(x, pos):
    """Formatter function to display y-axis ticks in millions (M) or thousands
       ↪(K)."""
    if x >= 1e9:
        return f'{x / 1e9:.1f}B'
    elif x >= 1e6:
        return f'{x / 1e6:.0f}M'
    return f'{x:.0f}'

from matplotlib.ticker import FuncFormatter

# Set up formatter
billion_formatter = FuncFormatter(format_billion)
million_formatter = FuncFormatter(format_million)
```

```
[24]: # Top 10 Directors by Total Profit ---
top_10_directors = director_total_profit.head(10)

plt.figure(figsize=(12, 6))
ax = sns.barplot(x='director', y='profit', data=top_10_directors)

# Apply formatter for billions/millions

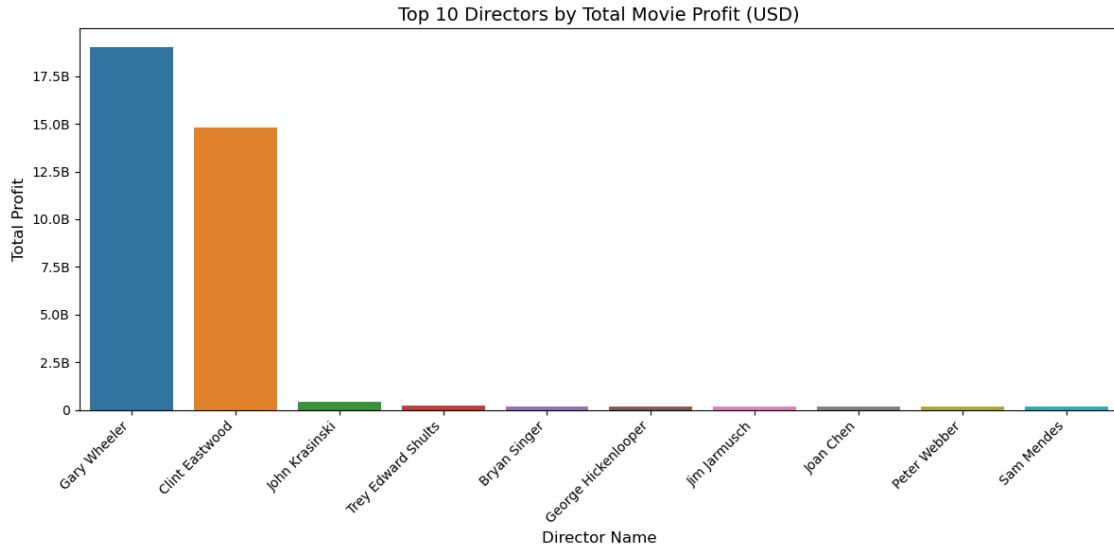
ax.yaxis.set_major_formatter(billion_formatter)

# Labels and formatting
plt.title('Top 10 Directors by Total Movie Profit (USD)', fontsize=14)
plt.xlabel('Director Name', fontsize=12)
plt.ylabel('Total Profit', fontsize=12)
```

```

plt.xticks(rotation=45, ha='right') # Rotate names to prevent overlap
plt.tight_layout()
plt.show()

```



```

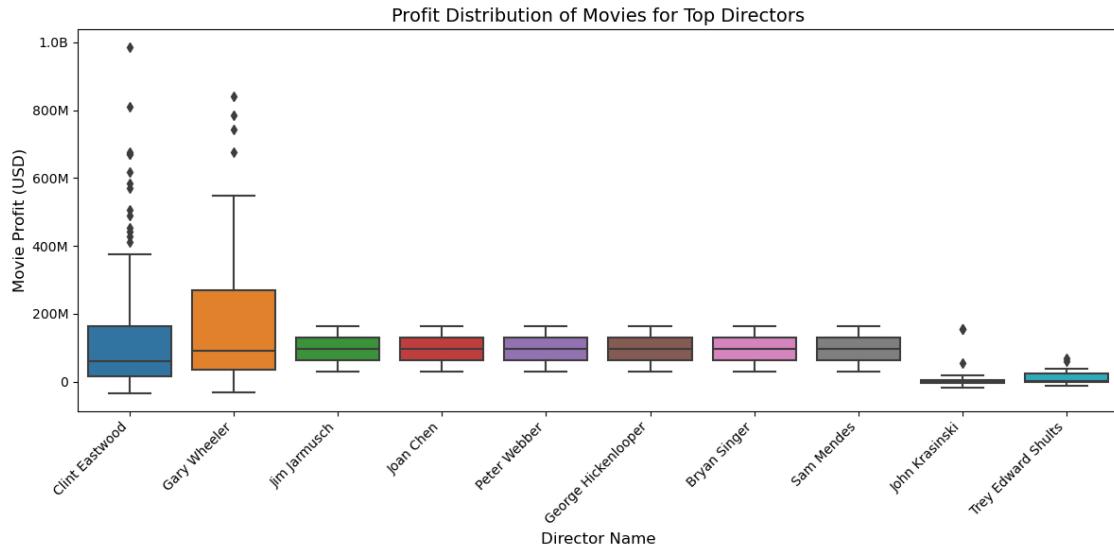
[25]: # Filter data for Top 10 Directors
top_directors_list = top_10_directors['director'].tolist()
filtered_df = merged_movie_info[merged_movie_info['director'].
                                isin(top_directors_list)]

# Create boxplot (shows distribution/risk of profits per director)
plt.figure(figsize=(12, 6))
ax = sns.boxplot(x='director', y='profit', data=filtered_df)

# Apply formatter for billions/millions
ax.yaxis.set_major_formatter(billion_formatter)

# Labels and formatting
plt.title('Profit Distribution of Movies for Top Directors', fontsize=14)
plt.xlabel('Director Name', fontsize=12)
plt.ylabel('Movie Profit (USD)', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

```



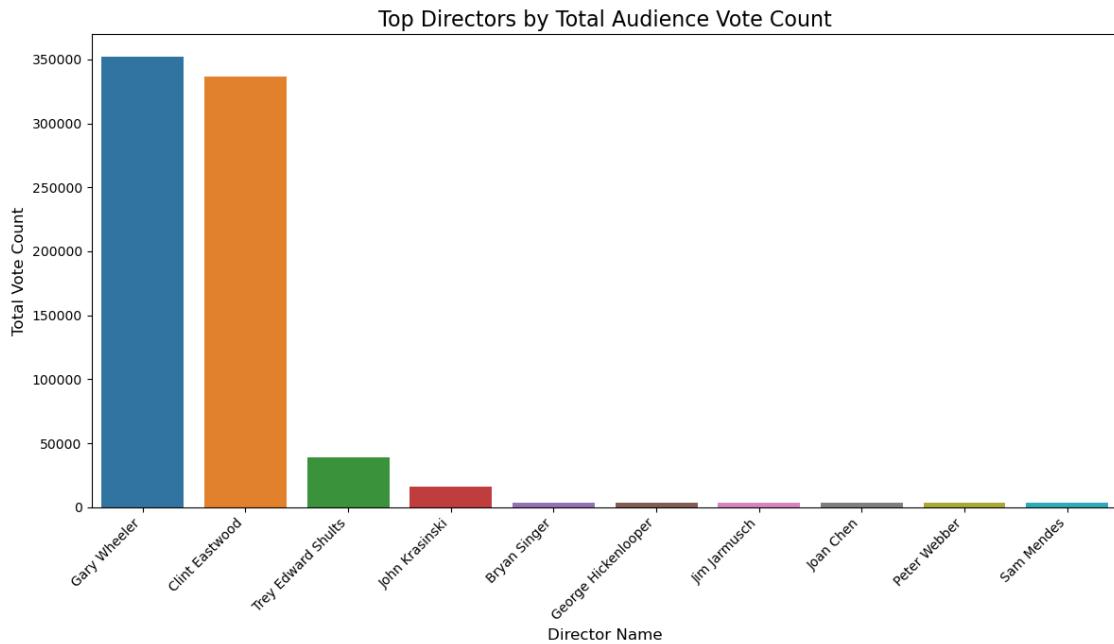
```
[26]: # Group by director and sum the total vote count and sort
director_total_votes = merged_movie_info.groupby('director')['vote_count'].
    sum().sort_values(ascending=False).reset_index()
director_total_votes.columns = ['Director Name', 'Total Vote Count']

# Select the top directors
top_directors_votes = director_total_votes.head(10)

# Bar Chart for Top Directors by Total Vote Count
plt.figure(figsize=(12, 7))

sns.barplot(x='Director Name', y='Total Vote Count', data=top_directors_votes)

plt.title('Top Directors by Total Audience Vote Count', fontsize=16)
plt.xlabel('Director Name', fontsize=12)
plt.ylabel('Total Vote Count', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
```

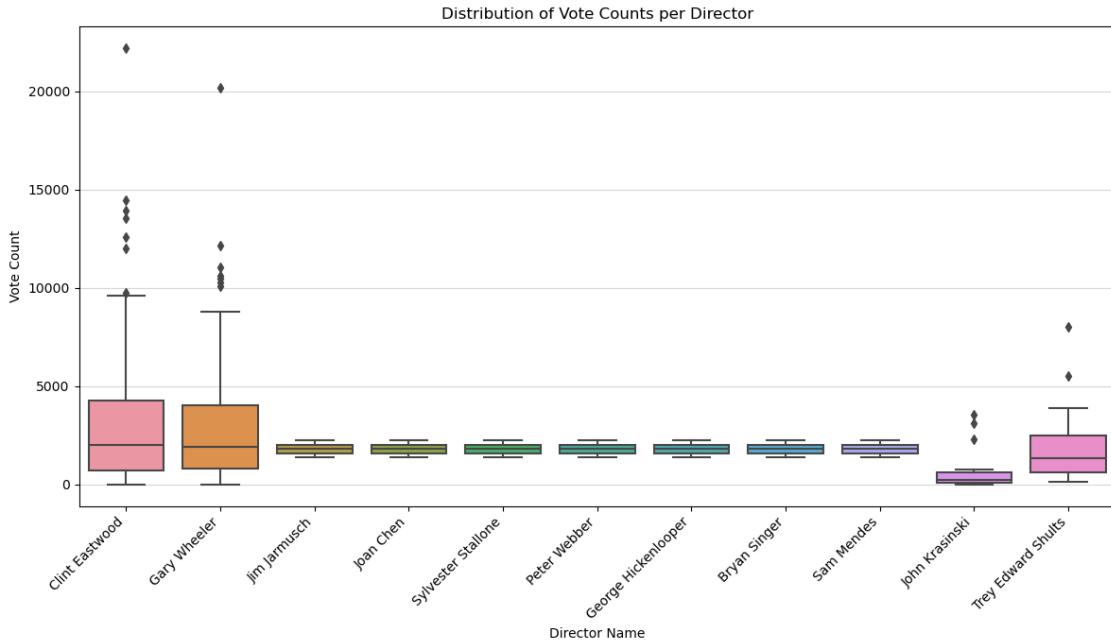


```
[27]: # Boxplot on distribution of vote counts per director
```

```
plt.figure(figsize=(12, 7))

sns.boxplot(x='director', y='vote_count', data=merged_movie_info)

plt.title('Distribution of Vote Counts per Director')
plt.xlabel('Director Name')
plt.ylabel('Vote Count')
plt.xticks(rotation=45, ha='right')
plt.grid(axis='y', alpha=0.5)
plt.tight_layout()
```



```
[28]: # Distribution of Top Writers by Total Movie Profits
# Split the combined writer names into individual rows
writers_split = merged_movie_info.assign(
    writer=merged_movie_info['writer'].str.split('|')
).explode('writer')

# Clean up whitespace or missing entries
writers_split['writer'] = writers_split['writer'].str.strip()
writers_split = writers_split.dropna(subset=['writer'])

writer_profit = (
    writers_split.groupby('writer')['profit']
    .sum()
    .sort_values(ascending=False)
    .head(10)
    .reset_index()
)

# Create bar plot
plt.figure(figsize=(10,6))
ax = plt.bar(writer_profit['writer'], writer_profit['profit'], color='skyblue')

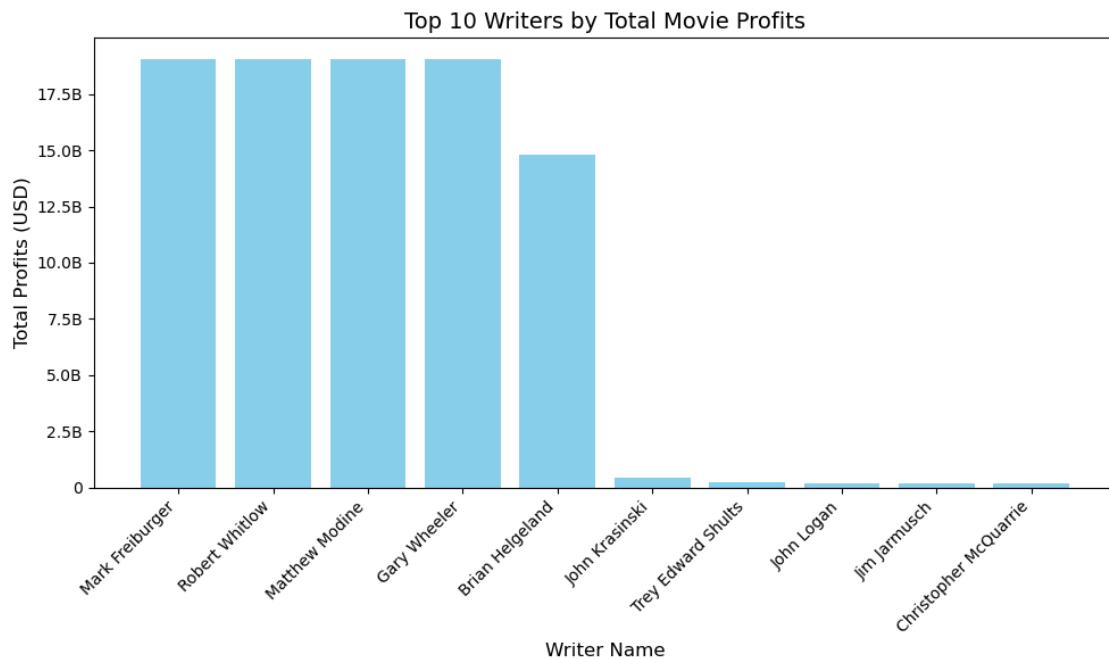
# Apply billions/millions formatter
plt.gca().yaxis.set_major_formatter(billion_formatter)

# Labels and formatting
```

```

plt.title('Top 10 Writers by Total Movie Profits', fontsize=14)
plt.xlabel('Writer Name', fontsize=12)
plt.ylabel('Total Profits (USD)', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

```



```

[29]: # Check Top 10 Writers
top_10_writers = (
    writers_split.groupby('writer')['profit']
    .sum()
    .sort_values(ascending=False)
    .head(10)
    .index.tolist()
)

# Include only movies from the Top 10 Writers
top_writers_data = writers_split[writers_split['writer'].isin(top_10_writers)]

# Create the boxplot
plt.figure(figsize=(12, 8))
ax = sns.boxplot(x='writer', y='profit', data=top_writers_data,
                  order=top_10_writers)

# Apply billions/millions formatter

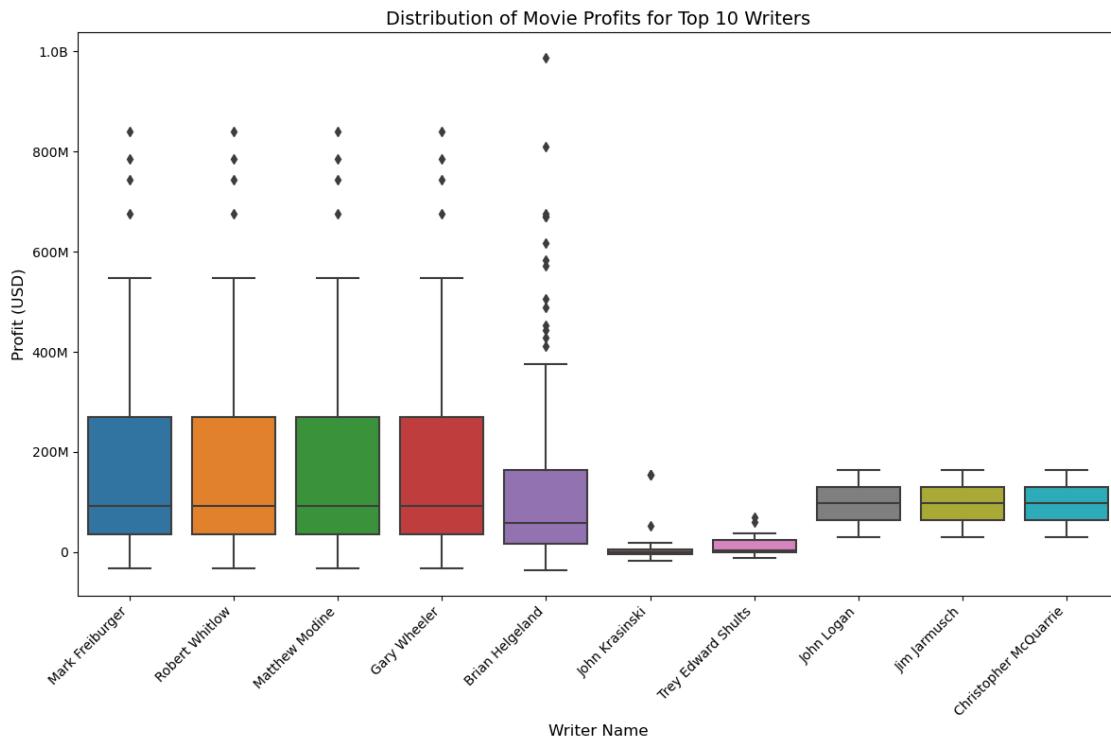
```

```

ax.yaxis.set_major_formatter(billion_formatter)

# Labels and formatting
plt.title('Distribution of Movie Profits for Top 10 Writers', fontsize=14)
plt.xlabel('Writer Name', fontsize=12)
plt.ylabel('Profit (USD)', fontsize=12)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

```



```

[30]: # writer vs vote counts
writer_votes = (
    writers_split.groupby('writer')['vote_count']
    .sum()
    .sort_values(ascending=False)
    .head(10)
    .reset_index()

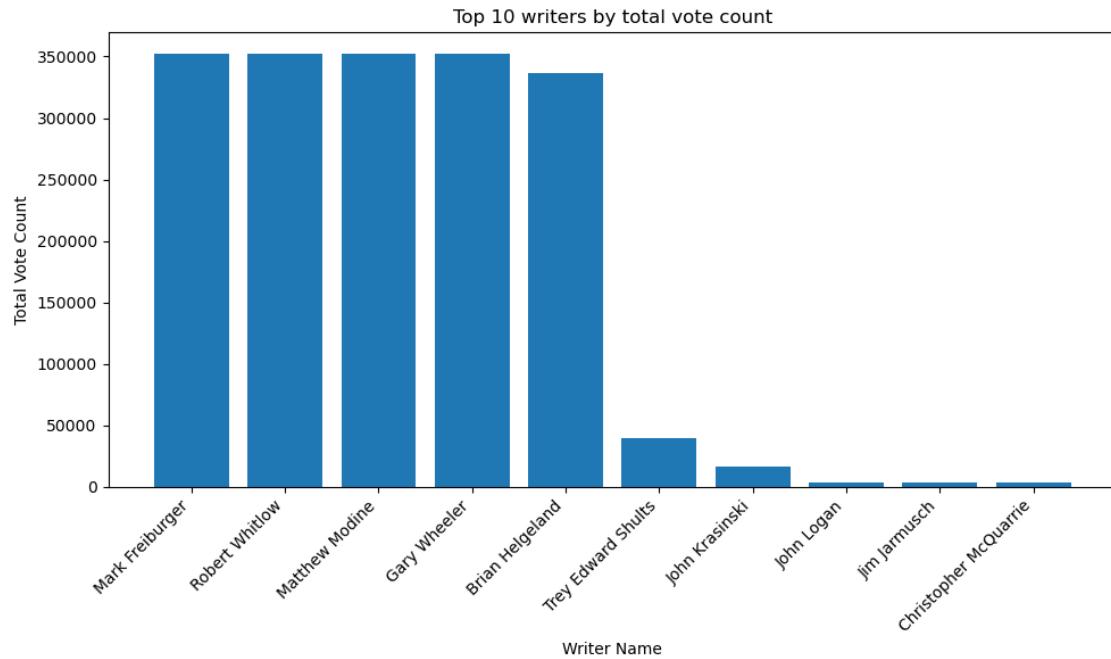
)
# plot
plt.figure(figsize=(10,6))
plt.bar(writer_votes['writer'],writer_votes['vote_count'])
plt.title('Top 10 writers by total vote count')

```

```

plt.xlabel('Writer Name')
plt.ylabel('Total Vote Count')
plt.xticks(rotation=45,ha='right')
plt.tight_layout()
plt.show()

```



```

[31]: # create notebook
top_10_writers_by_votes = (writers_split.groupby('writer')['vote_count']
                           .sum()
                           .sort_values(ascending=False)
                           .head(10)
                           .index.tolist()
                           )

# Include only movies from the top 10 writers
# This provides the distribution of individual vote counts for the boxplot
top_writers_vote_data = writers_split[writers_split['writer'].isin(top_10_writers_by_votes)]

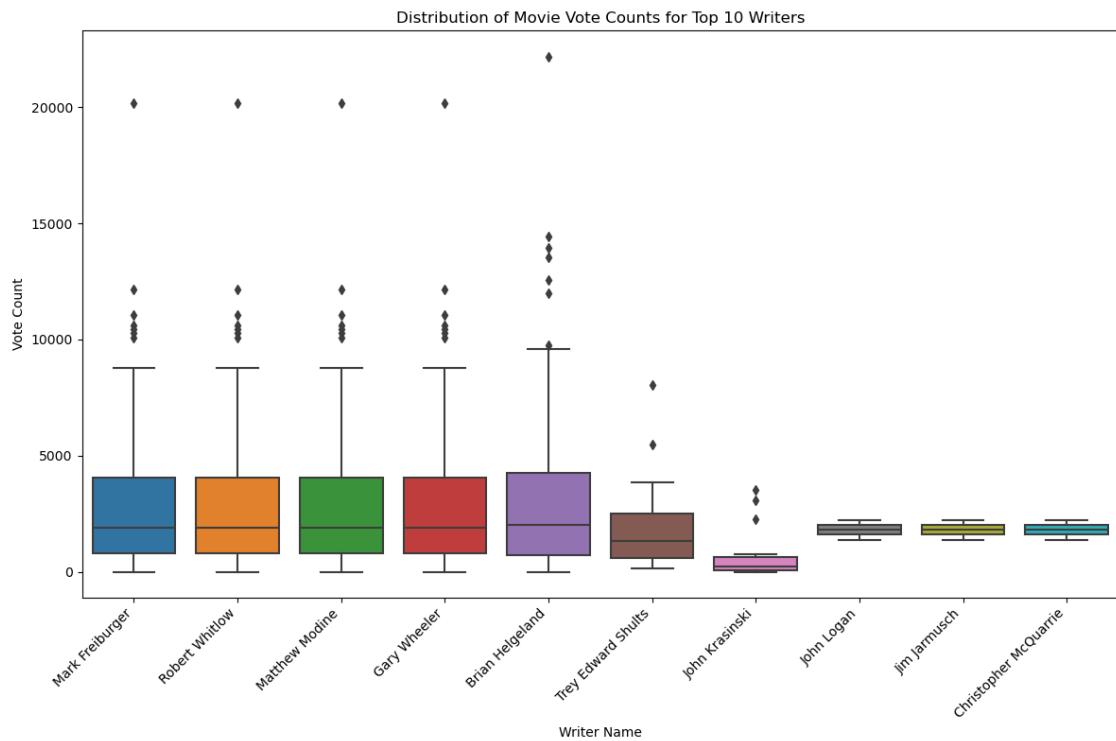
# Create the boxplot
plt.figure(figsize=(12, 8))
# Use 'sns.boxplot' with 'order' to display the writers in descending order of total votes
sns.boxplot(x='writer', y='vote_count', data=top_writers_vote_data, order=top_10_writers_by_votes)

```

```

plt.title('Distribution of Movie Vote Counts for Top 10 Writers')
plt.xlabel('Writer Name')
plt.ylabel('Vote Count')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

```



0.6 SECTION 5: MOVIE GENRE ANALYSIS BY RATINGS AND PROFITS

This section shows data preparation and Analysis to establish which genres have the highest ratings & to assess which genres generate the highest gross revenue and profit to give our stakeholders an idea on which genres they can focus on when they start creating movies.

```
[32]: # check for null values in the dataset
df.isna().sum()
```

```
[32]: movie_titles          1
genres                  146
original_language       157
release_date            1
studio                 44280
production_budget       1
domestic_gross          44279
```

```

foreign_gross          44475
worldwide_gross        44279
revenue                  1
runtime                 394
popularity                146
production_companies      146
production_countries      146
vote_average                146
vote_count                  146
dtype: int64

```

```
[33]: # Since revenue is the same as gross we will drop the domestic,
# foreign and worldwide gross columns since they have a lot of
# missing values and the revenue column has less missing values
# Drop all the columns not needed in this analysis.
columns_to_drop = [
    'domestic_gross',
    'foreign_gross',
    'worldwide_gross',
    'studio',
    'original_language',
]
df = df.drop(columns=columns_to_drop)
```

```
[34]: # Convert all values to string, then remove $ and comma, and finally convert to numeric
def clean_currency_column(series):
    cleaned_series = series.astype(str).str.replace('$', '', regex=False).str.replace(',', '', regex=False)
    return pd.to_numeric(cleaned_series, errors='coerce')

# Parse JSON-like string in the 'genres' column
def extract_genre_names(genres_str):
    if pd.isna(genres_str) or genres_str in ('[]', ''):
        return []
    try:
        # Use simple string replacement to fix single quotes to double quotes
        for valid_JSON:
            genres_str_fixed = genres_str.replace("''", "''")
            genres_list = json.loads(genres_str_fixed)
            return [g['name'] for g in genres_list]
    except (TypeError, json.JSONDecodeError):
        return []
```

```
[35]: 
```

```

df['revenue'] = clean_currency_column(df['revenue']) # Worldwide revenue
df['production_budget'] = clean_currency_column(df['production_budget']) # ↵
    ↵Production budget
df['vote_average'] = pd.to_numeric(df['vote_average'], errors='coerce') # ↵
    ↵Ratings

# Calculate Profit
df['profit'] = df['revenue'] - df['production_budget']

# Drop rows where key financial data is missing or zero (not reliable for ↵
    ↵profit/revenue analysis)
df.dropna(subset=['revenue', 'production_budget', 'profit', 'vote_average'], ↵
    ↵inplace=True)
df = df[(df['revenue'] > 1e4) & (df['production_budget'] > 1e4)] # Filter out ↵
    ↵near-zero values

# Prepare Genres (Exploding the DataFrame)
df['Genres'] = df['genres'].apply(extract_genre_names)
# Filter out movies with no valid genre
df_genre = df[df['Genres'].apply(lambda x: len(x) > 0)].copy()
# Explode the DataFrame to have one row per movie-genre combination
df_exploded = df_genre.explode('Genres')

```

[36]:

```

# Season function so as to work with the seasons instead of months
df_exploded['release_date'] = pd.to_datetime(df_exploded['release_date'], ↵
    ↵errors='coerce')
df_exploded.dropna(subset=['release_date'], inplace=True) # drop null

def get_season(month):
    """Maps month number to a cinematic release season."""
    if month in [12, 1, 2]: # December, January, February
        return 'Winter'
    elif month in [3, 4, 5]: # March, April, May
        return 'Spring'
    elif month in [6, 7, 8]: # June, July, August
        return 'Summer'
    else: # September, October, November
        return 'Fall'

df_exploded['Season'] = df_exploded['release_date'].dt.month.apply(get_season)

```

[37]:

```

sns.set_theme(style="whitegrid") # so as to have a standard theme in the plots

# 1. Genre vs Profit (Total and Average)
# Total Profit

```

```

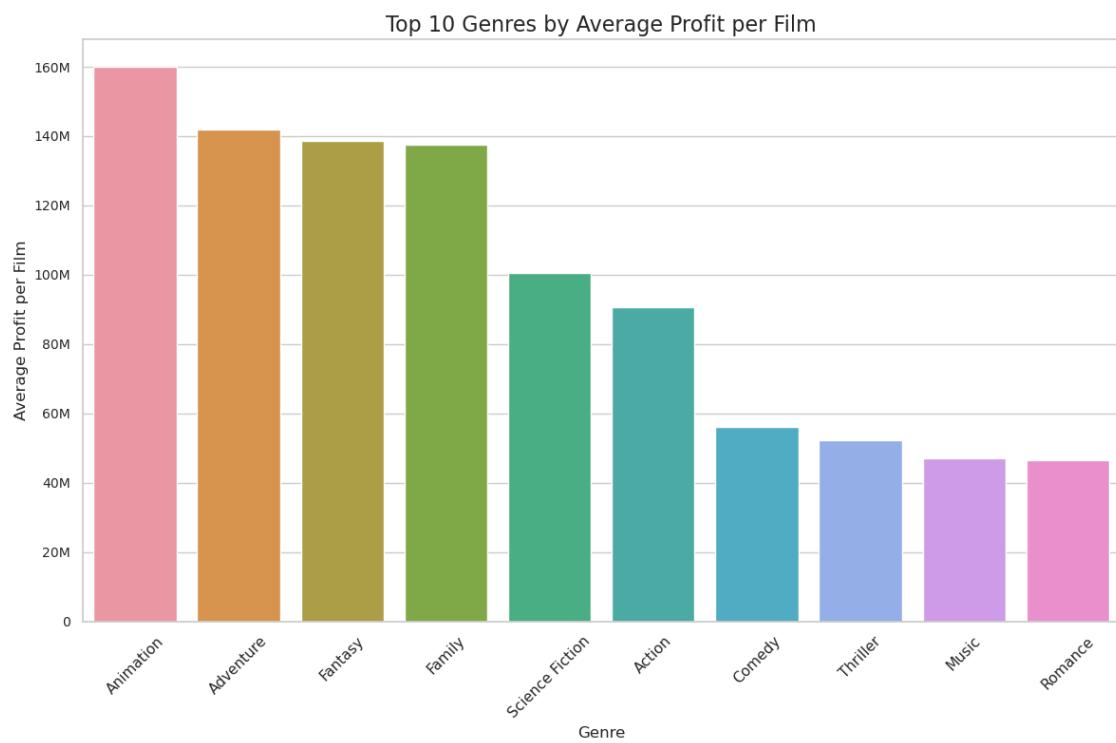
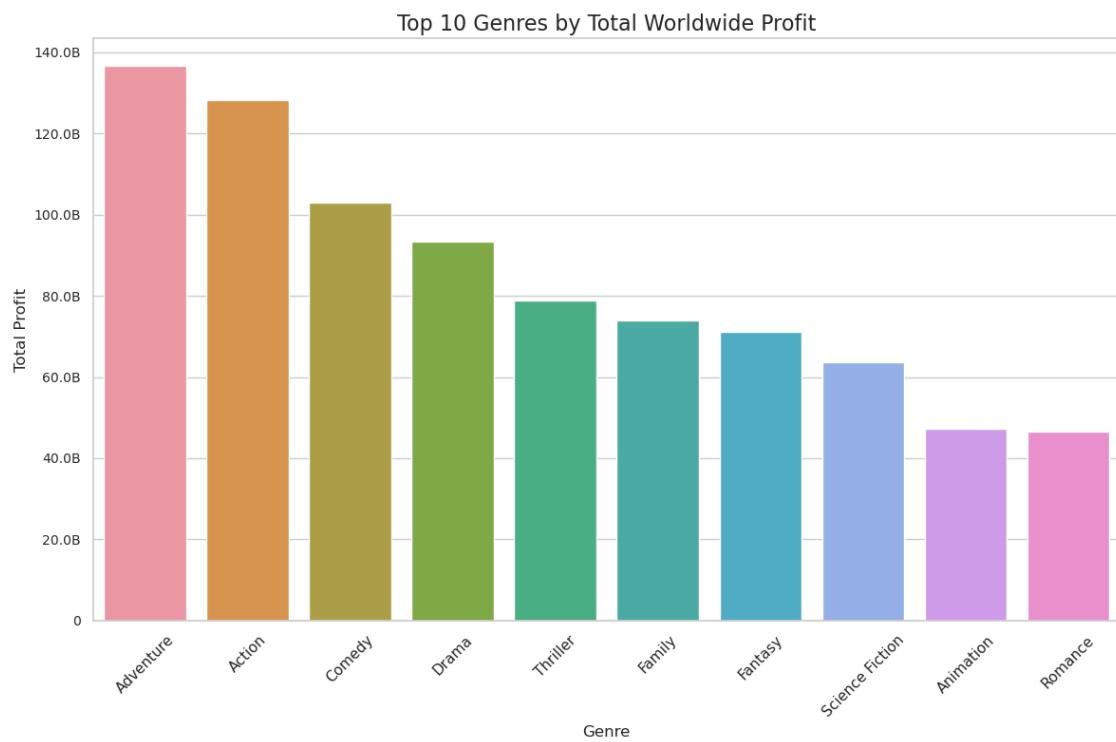
genre_total_profit = df_exploded.groupby('Genres')['profit'].sum() .
    ↪sort_values(ascending=False).head(10)

# Average Profit
genre_avg_profit = df_exploded.groupby('Genres')['profit'].mean() .
    ↪sort_values(ascending=False).head(10)

# Plot1. Total Profit
plt.figure(figsize=(12,8))
sns.set_palette("crest")
sns.barplot(x=genre_total_profit.index, y=genre_total_profit.values)
plt.title('Top 10 Genres by Total Worldwide Profit', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Total Profit')
plt.gca().yaxis.set_major_formatter(billion_formatter)
plt.xticks(rotation=45)
plt.yticks(fontsize=10)
plt.tight_layout()
# plt.savefig('images/genre_total_profit.png')
plt.show()

# Plot 2: Average Profit
plt.figure(figsize=(12,8))
sns.set_palette("crest")
sns.barplot(x=genre_avg_profit.index, y=genre_avg_profit.values)
plt.title('Top 10 Genres by Average Profit per Film', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Average Profit per Film')
plt.gca().yaxis.set_major_formatter(million_formatter)
plt.xticks(rotation=45)
plt.yticks(fontsize=10)
plt.tight_layout()
# plt.savefig('images/genre_avg_profit.png')
plt.show()

```



```
[38]: # 2 Genre vs Revenue (Total and Average)

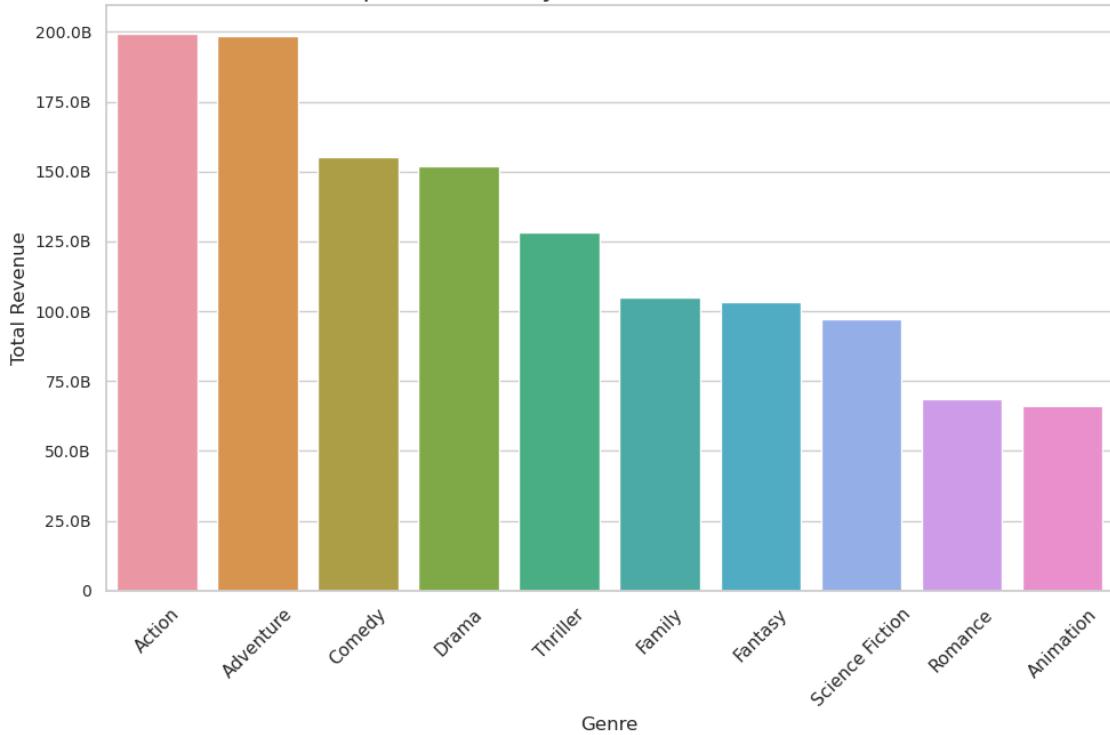
# Total Revenue
genre_total_revenue = df_exploded.groupby('Genres')['revenue'].sum().
    sort_values(ascending=False).head(10)

# Average Revenue
# Re-calculate average on unique movies per genre to avoid skewing the average
# (Although df_exploded.groupby().mean() is statistically okay, this is
# conceptually cleaner)
unique_titles = df_genre.explode('Genres')
genre_avg_revenue = unique_titles.groupby('Genres')['revenue'].mean().
    sort_values(ascending=False).head(10)

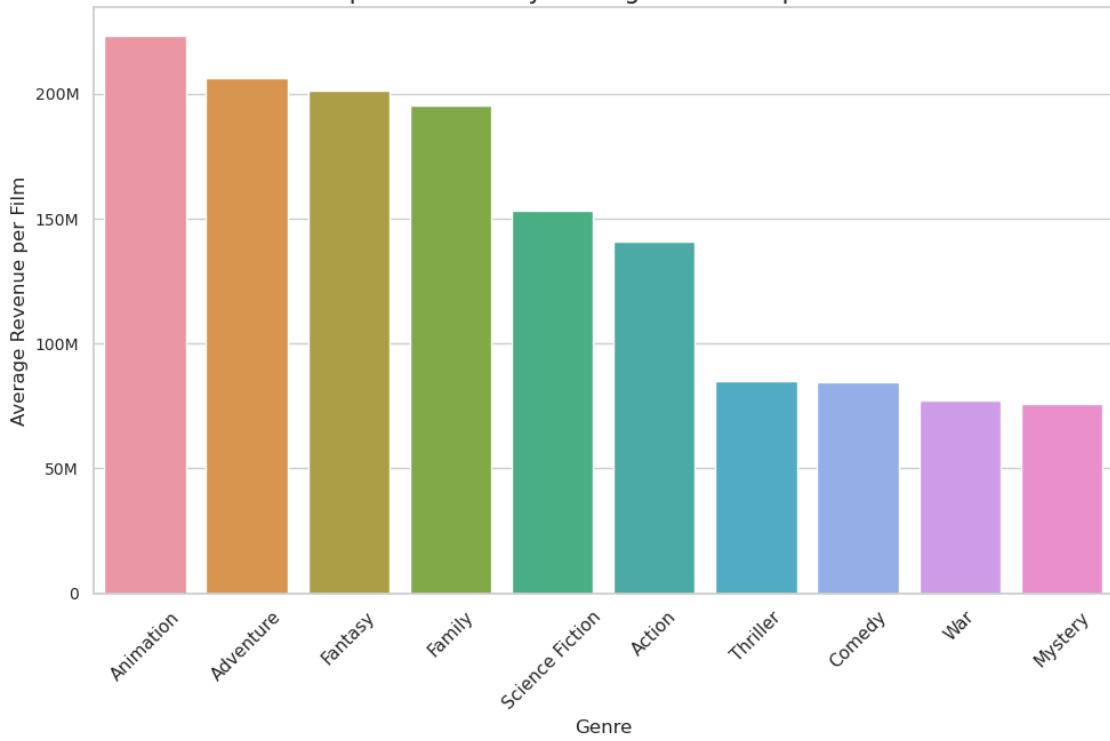
# Plot 1: Total Revenue
plt.figure(figsize=(10,7))
sns.set_palette("crest")
sns.barplot(x=genre_total_revenue.index, y=genre_total_revenue.values)
plt.title('Top 10 Genres by Total Worldwide Revenue', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Total Revenue')
plt.gca().yaxis.set_major_formatter(billion_formatter)
plt.xticks(rotation=45)
plt.yticks(fontsize=10)
plt.tight_layout()
# plt.savefig('images/genre_vs_total_revenue.png')
plt.show()

# Plot 2: Average Revenue
plt.figure(figsize=(10,7))
sns.set_palette("crest")
sns.barplot(x=genre_avg_revenue.index, y=genre_avg_revenue.values)
plt.title('Top 10 Genres by Average Revenue per Film', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Average Revenue per Film')
plt.gca().yaxis.set_major_formatter(million_formatter)
plt.xticks(rotation=45)
plt.yticks(fontsize=10)
plt.tight_layout()
# plt.savefig('images/genre_vs_revenue_avg.png')
plt.show()
```

Top 10 Genres by Total Worldwide Revenue



Top 10 Genres by Average Revenue per Film



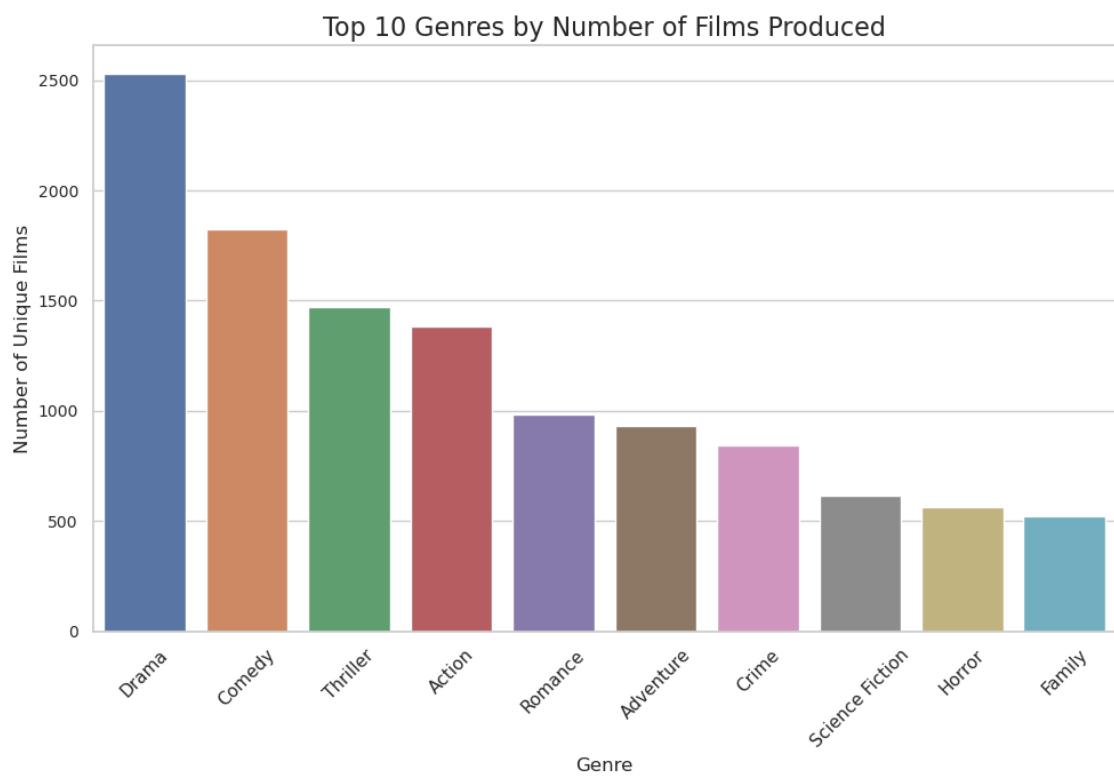
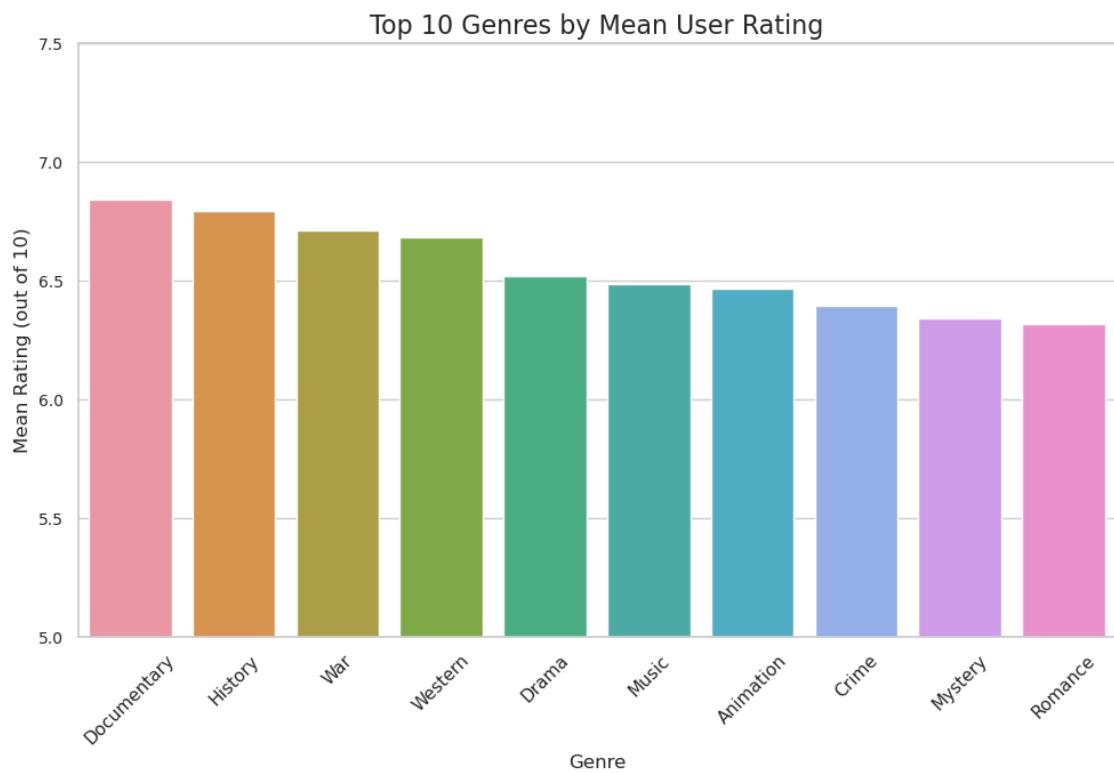
```
[39]: # 3 Genre vs Ratings (Mean Rating and Movie Count)

# Mean Rating
genre_mean_rating = df_exploded.groupby('Genres')['vote_average'].mean() .
    sort_values(ascending=False).head(10)

# Movie Count
genre_count = df_exploded.groupby('Genres')['movie_titles'].nunique() .
    sort_values(ascending=False).head(10)

# Plot 1: Mean Rating
plt.figure(figsize=(10, 7))
sns.set_palette("flare")
sns.barplot(x=genre_mean_rating.index, y=genre_mean_rating.values)
plt.title('Top 10 Genres by Mean User Rating', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Mean Rating (out of 10)')
plt.xticks(rotation=45)
plt.ylim(5, 7.5)
plt.yticks(fontsize=10)
plt.tight_layout()
# plt.savefig('images/genre_vs_ratings_count.png')
plt.show()

# Plot 2: Movie Count
plt.figure(figsize=(10, 7))
sns.set_palette("deep")
sns.barplot(x=genre_count.index, y=genre_count.values)
plt.title('Top 10 Genres by Number of Films Produced', fontsize=16)
plt.xlabel('Genre')
plt.ylabel('Number of Unique Films')
plt.xticks(rotation=45)
plt.yticks(fontsize=10)
plt.tight_layout()
# plt.savefig('images/genre_vs_movie_count.png')
plt.show()
```



0.7 SECTION 6: MOVIE PRODUCTION BUDGET VS. WORLDWIDE REVENUE

The scatter plot below illustrates the relationship between a movie's production budget and its worldwide revenue. The plot shows a positive trend: higher production budgets generally lead to higher worldwide revenues. However, the spread of points around the break-even line indicates that bigger budgets don't always guarantee profits, as some expensive films still underperform.

```
[40]: # 4 Budget vs Revenue
df_plot = df.copy()

plt.figure(figsize=(10, 8))
sns.set_palette("viridis")
sns.scatterplot(
    x='production_budget',
    y='revenue',
    data=df_plot,
    alpha=0.6,
    s=80,
    linewidth=0
)

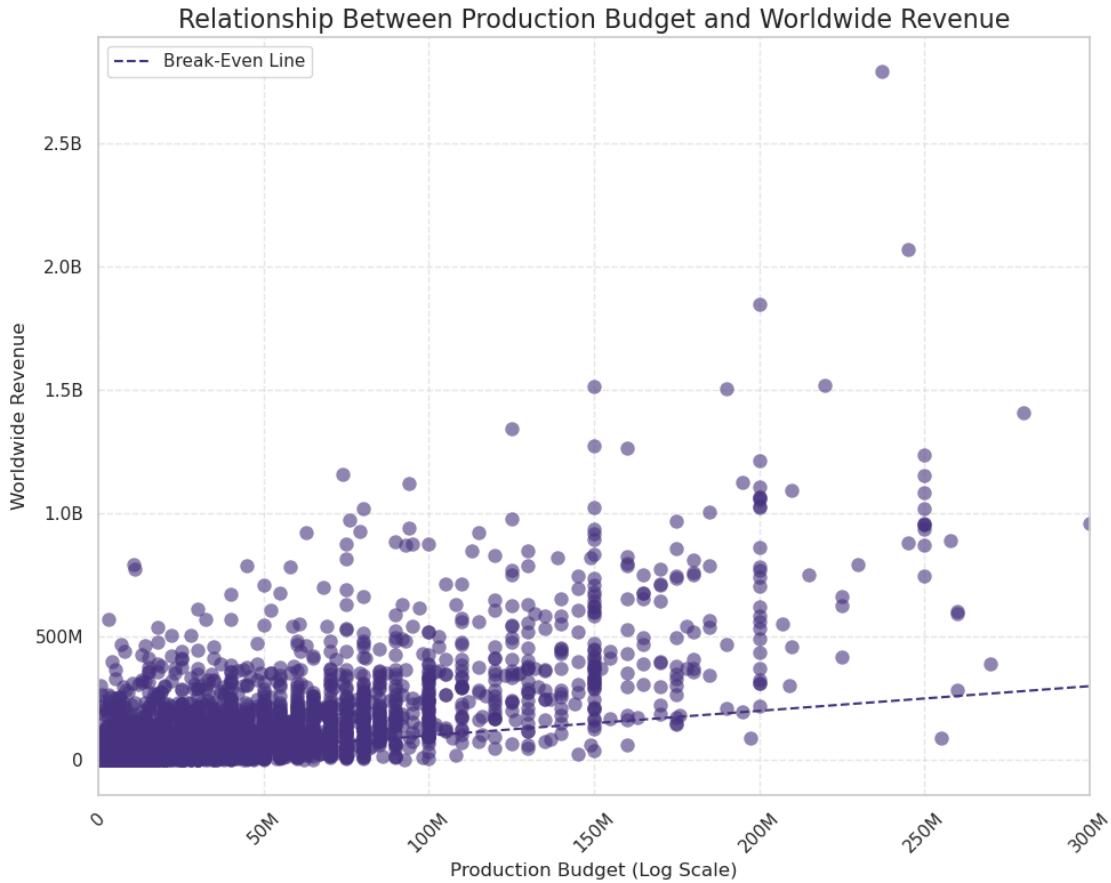
# Break-even line
max_value = max(df_plot['production_budget'].max(), df_plot['revenue'].max())
plt.plot([1, max_value], [1, max_value], linestyle='--', label='Break-Even Line')

plt.title('Relationship Between Production Budget and Worldwide Revenue', fontsize=16)
plt.xlabel('Production Budget (Log Scale)')
plt.ylabel('Worldwide Revenue')

# Apply custom formatters
plt.gca().xaxis.set_major_formatter(billion_formatter)
plt.gca().yaxis.set_major_formatter(billion_formatter)
plt.tick_params(axis='x', rotation=45)

plt.xlim(0, 300_000_000) # Show only budgets $300M

plt.grid(True, linestyle='--', alpha=0.5)
plt.legend()
plt.tight_layout()
# plt.savefig('images/budget_vs_revenue.png')
plt.show()
```



0.8 SECTION 7: RATINGS VS. PROFITABILITY BY BUDGET CATEGORY

The chart below shows that movies with higher audience ratings generally earn higher profits, especially in the high and blockbuster budget categories. Lower-budget films tend to cluster around smaller profits, while larger-budget films show greater variation indicating that bigger budgets increase both potential reward and financial risk.

[41]: # 5 Ratings vs Profits (Budget as Hue)

```
df_plot = df.copy()

# Create a categorical budget size column for color coding
df_plot['Budget_Category'] = pd.cut(
    df_plot['production_budget'],
    bins=[-np.inf, 1e7, 5e7, 1.5e8, np.inf],
    labels=['Low (<$10M)', 'Mid (<$50M)', 'High (<$150M)', 'Blockbuster'],
    )

```

```

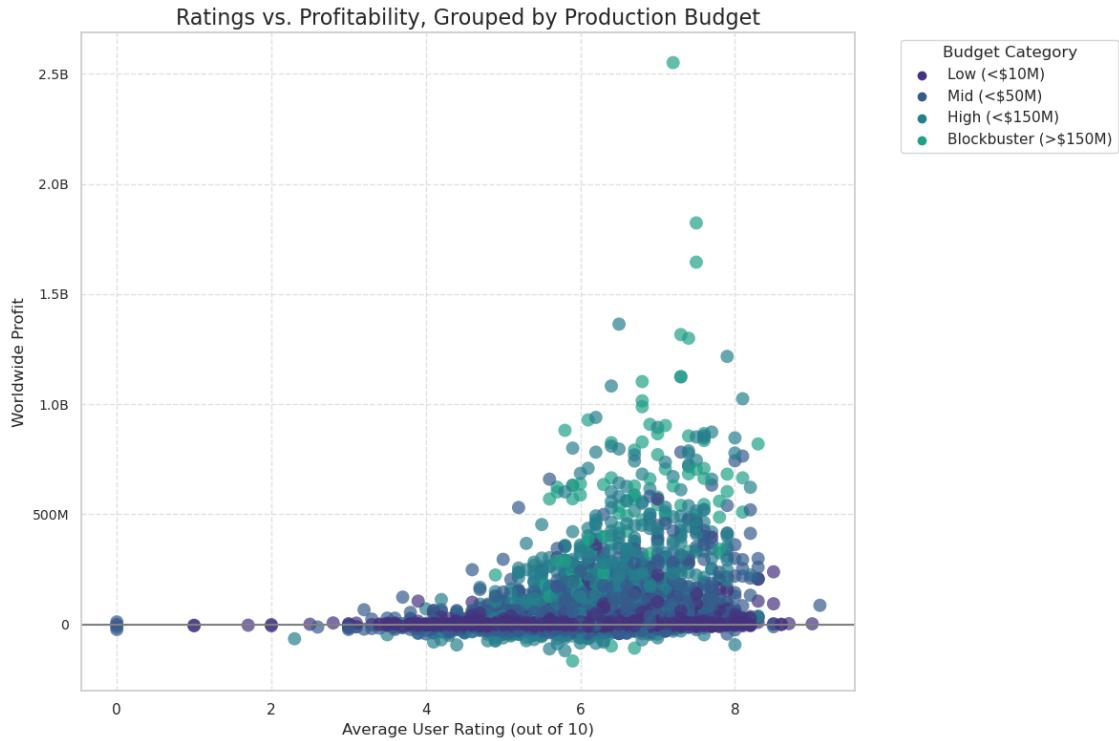
# Scatter plot with Rating on X, Profit on Y, and Budget Category as Hue
plt.figure(figsize=(12, 8))
sns.set_palette("viridis")
sns.scatterplot(
    x='vote_average',
    y='profit',
    hue='Budget_Category',
    data=df_plot,
    alpha=0.7,
    s=100,
    linewidth=0,
    hue_order=['Low (<$10M)', 'Mid (<$50M)', 'High (<$150M)', 'Blockbuster' ↴(>$150M)']
)
# Draw a line at Profit = 0
plt.axhline(0, color='gray', linestyle='--')

plt.title('Ratings vs. Profitability, Grouped by Production Budget', ↴
          fontsize=16)
plt.xlabel('Average User Rating (out of 10)')
plt.ylabel('Worldwide Profit')

# Apply custom formatter to Y-axis (Profit)
plt.gca().yaxis.set_major_formatter(billion_formatter)
plt.tick_params(axis='y', labelsize=10)

plt.legend(title='Budget Category', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
#plt.savefig('images/ratings_vs_profits_budget.png', dpi=300)
plt.show()

```



0.9 SECTION 8: PROFITABILITY BY MOVIE RELEASE SEASON

Movies released in summer and spring tend to achieve the highest profits, likely due to school holidays and major blockbuster releases. In contrast, fall and winter films show lower average profits, suggesting that release timing plays a key role in maximizing box office returns.

[42]: # 6 Seasons vs Profits

```
season_order = ['Winter', 'Spring', 'Summer', 'Fall']

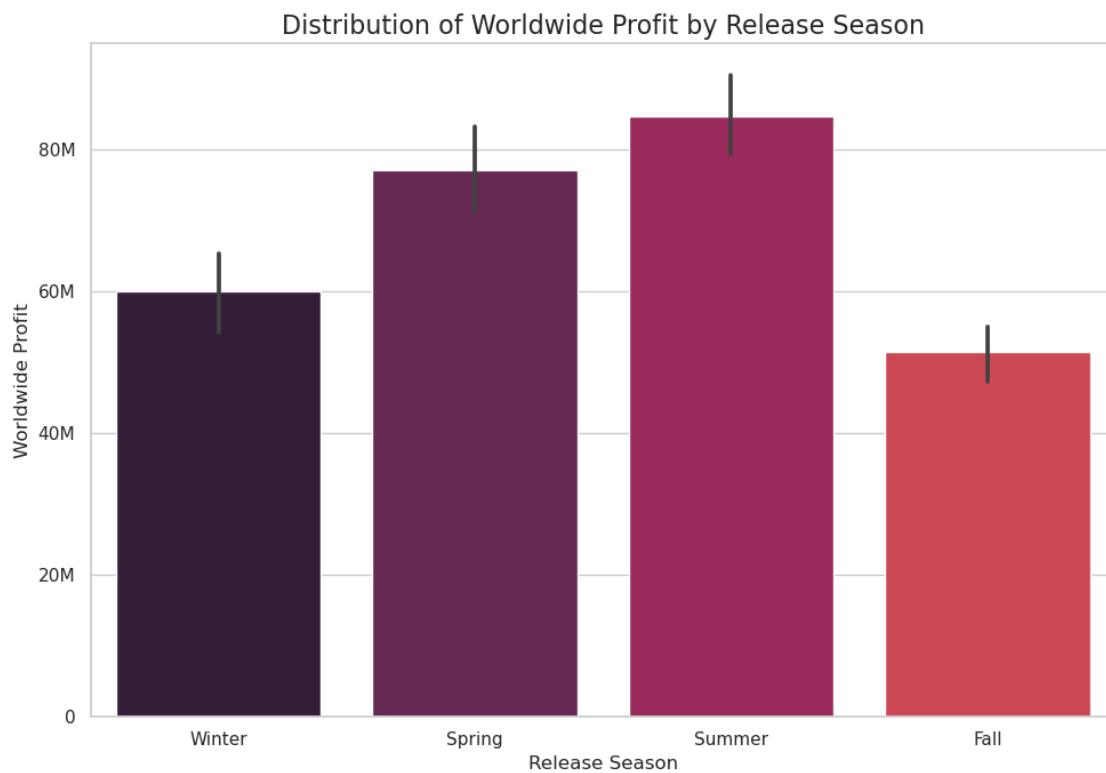
plt.figure(figsize=(10, 7))
sns.set_palette("rocket")
sns.barplot(
    x='Season',
    y='profit',
    data=df_exploded,
    order=season_order,
)
plt.title('Distribution of Worldwide Profit by Release Season', fontsize=16)
plt.xlabel('Release Season')
plt.ylabel('Worldwide Profit')
```

```

plt.gca().yaxis.set_major_formatter(billion_formatter)

plt.tight_layout()
# plt.savefig('images/seasons_vs_profits.png')
plt.show()

```



0.10 SECTION 9: CONCLUSIONS AND RECOMMENDATIONS

0.10.1 CONCLUSIONS

- Movie length vs audience ratings:** The longer movies tend to be rated more favorably
- Directors and writers Vs Movie Performance:** The directors generating high revenue to be considered for better returns.
- Genres effects on ratings, gross revenue and profit:** Genres with high total profit are better options for studios aiming for consistent returns
- Movie release season vs profit :** Studios to prioritize Summer and Spring for major release to maximize revenue

0.10.2 RECOMMENDATIONS

- Movie length vs audience ratings:** Content creators to consider longer movie run times, they probably give an allowance for deeper story telling.
- Directors and writers Vs Movie Performance:** The studio to consider prioritizing di-

rectors who have released movies that generated high revenues e.g Clint Eastwood and Gary Wheeler

3. **Genres effects on ratings, gross revenue and profit:** The studios to consider prioritizing genres with high total profit for consistent returns
4. **Movie release season vs profit :** Seasonality should be considered in release planning and marketing strategies.